

Statistical Learning Group 5

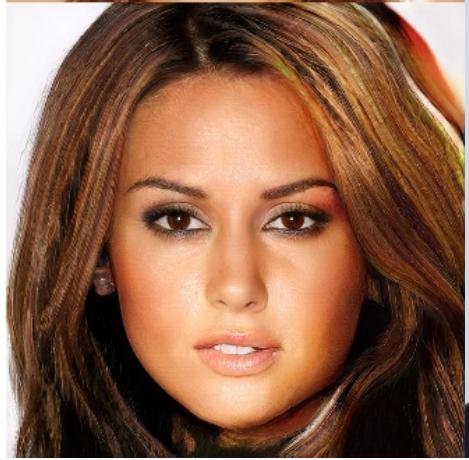
Generative Models

Anup Tuladhar

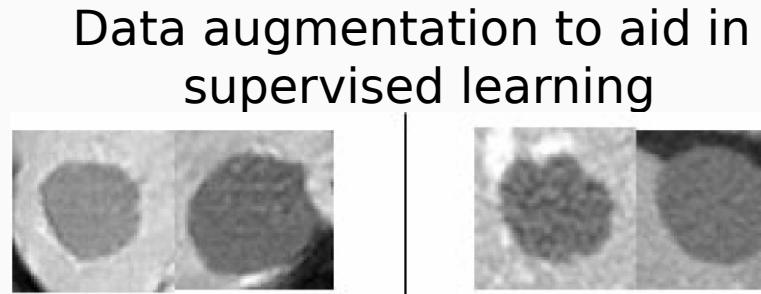
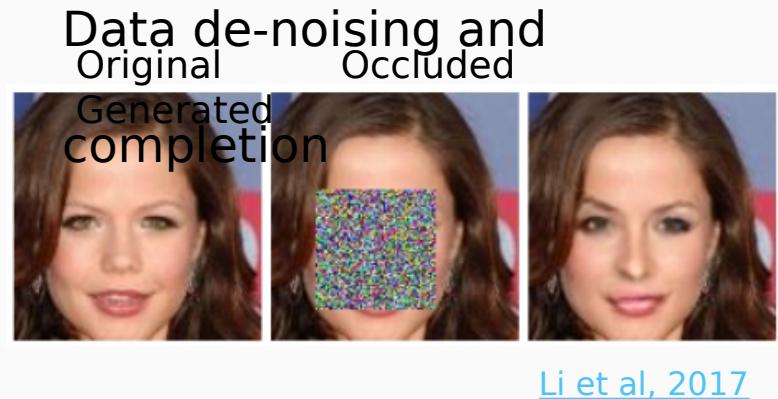
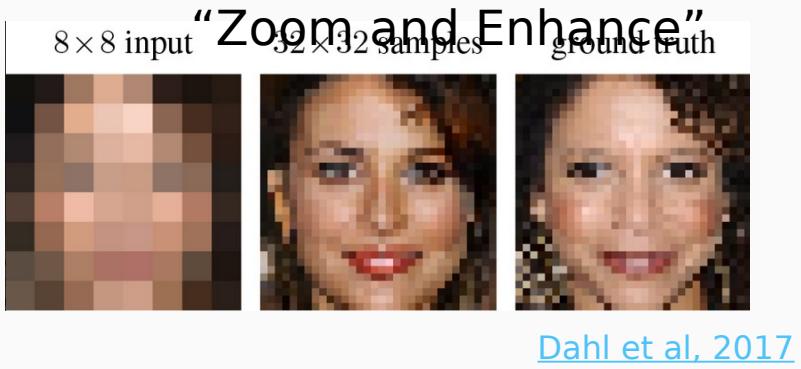
May 31st, 2019



Which reality TV star is fake?



What are the interesting use cases for generative models

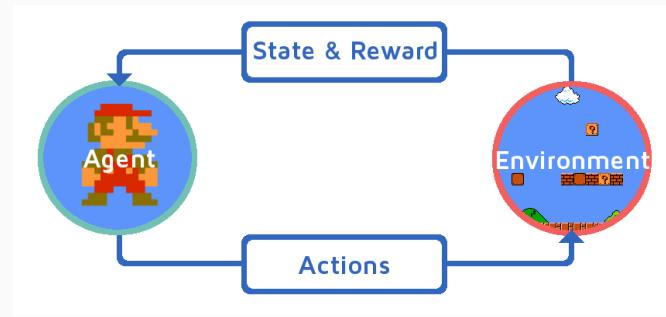


Machine learning categories

Reinforcement Learning

Learn a **policy** for an agent to maximize its **reward** in an environment

E.g. Q-learning, SARSA



Machine learning categories

Reinforcement Learning

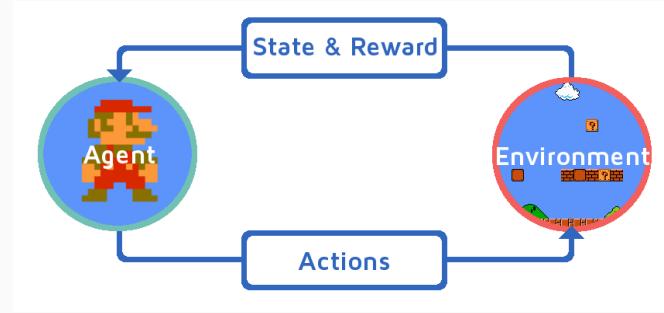
Learn a **policy** for an agent to maximize its **reward** in an environment

E.g. Q-learning, SARSA

Supervised Learning

Learn a model that knows how to map **data** to **label**

E.g. Decision trees, Random forests, Support vector machines



Machine learning categories

Reinforcement Learning

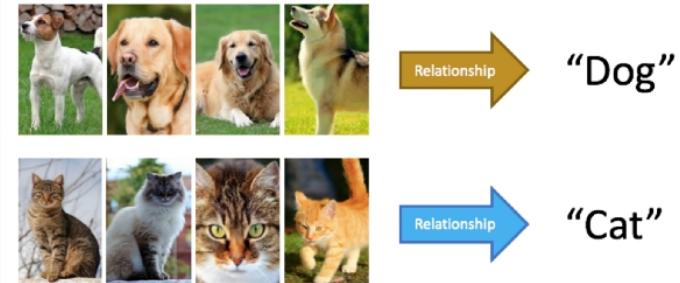
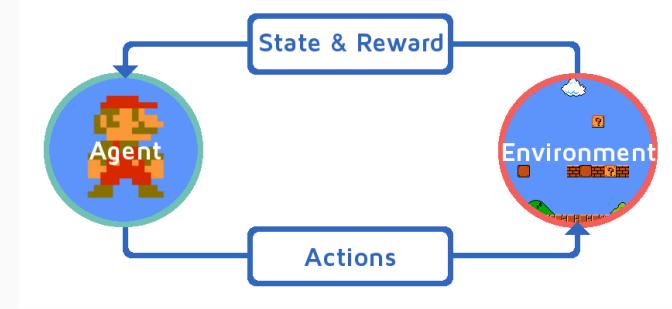
Learn a **policy** for an agent to maximize its **reward** in an environment

E.g. Q-learning, SARSA

Supervised Learning

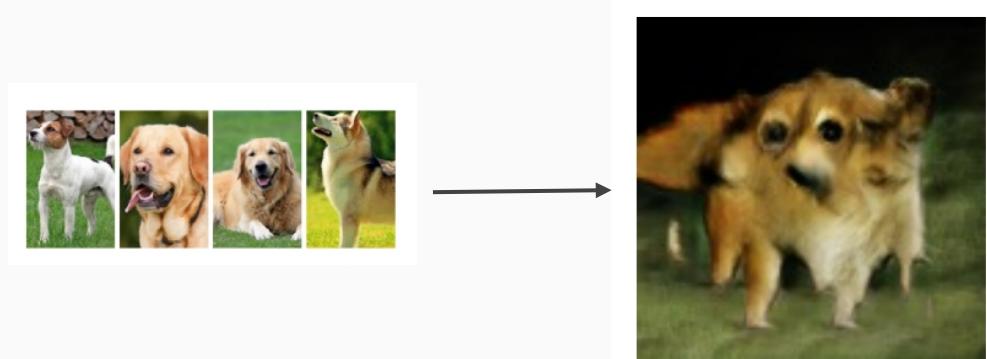
Learn a model that knows how to map **data** to **label**

E.g. Decision trees, Random forests, Support vector machines



Unsupervised Learning

Learn the underlying structure of the **data** in the absence of labels

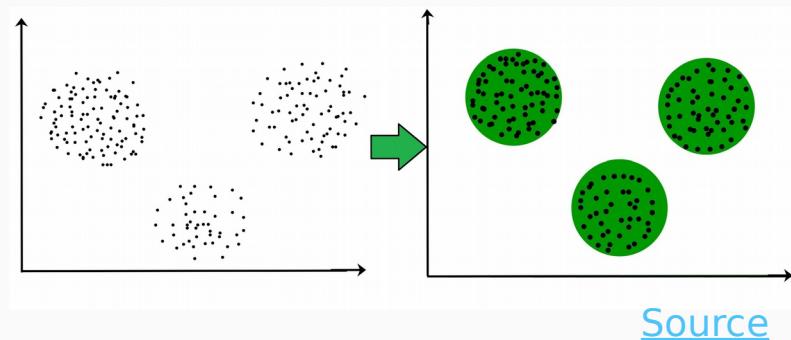


Types of Unsupervised Learning

Clustering

Group the data into based on the similarity, or dissimilarity, of their features

E.g. k-means clustering, hierarchical clustering

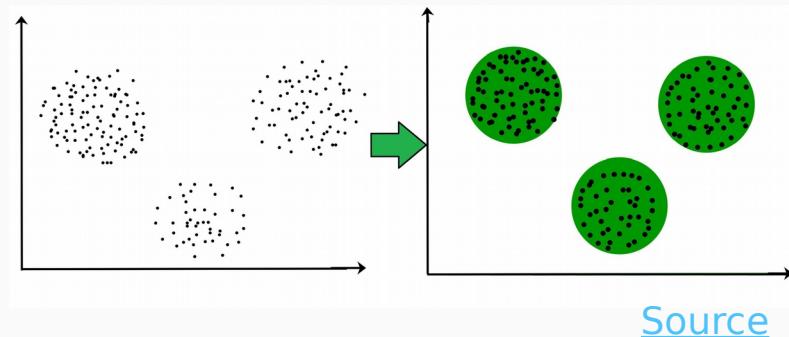


Types of Unsupervised Learning

Clustering

Group the data into based on the similarity, or dissimilarity, of their features

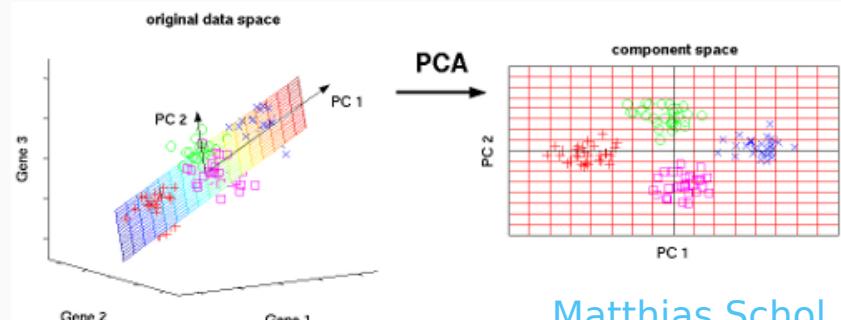
E.g. k-means clustering, hierarchical clustering



Dimensionality reduction

Distill the feature space of the data down to a **smaller feature space** that still captures the **most important sources of variation**

E.g. Principal component analysis, autoencoders



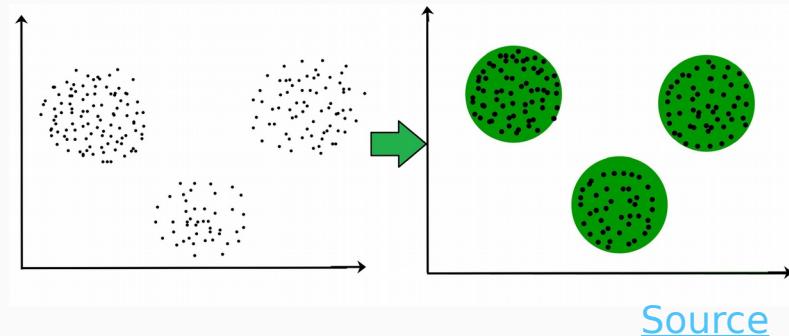
[Matthias Scholz](#)

Types of Unsupervised Learning

Clustering

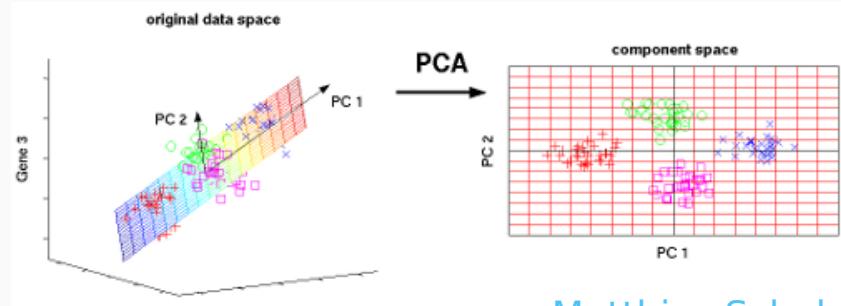
Group the data into based on the similarity, or dissimilarity, of their features

E.g. k-means clustering, hierarchical clustering



Dimensionality reduction

Distill the feature space of the data down to a **smaller feature space** that still captures the **most important sources of variation**

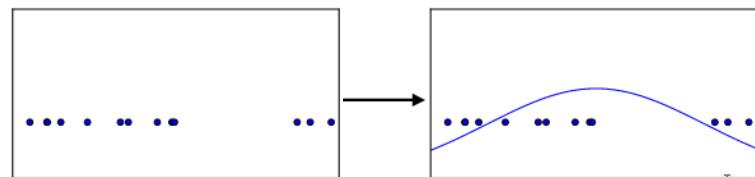


[Matthias Scholz](#)

E.g. Principal component analysis, autoencoders

Density estimation

Determine the underlying distribution of the data



[Ian Goodfellow](#)

Generative Modeling

General Goal: Model the joint data distribution
 $p(\text{observed}, \text{target})$

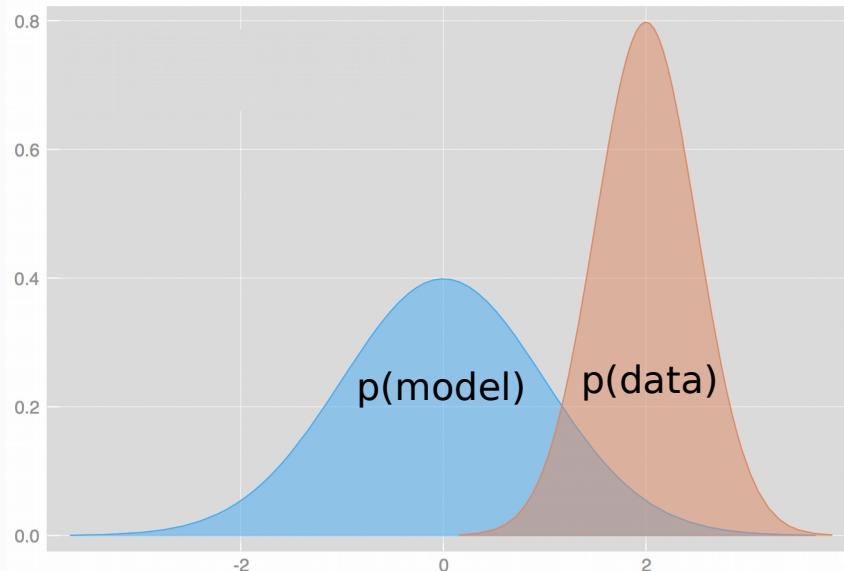
Generative Modeling

**General Goal: Model the joint data distribution
 $p(\text{observed, target})$**

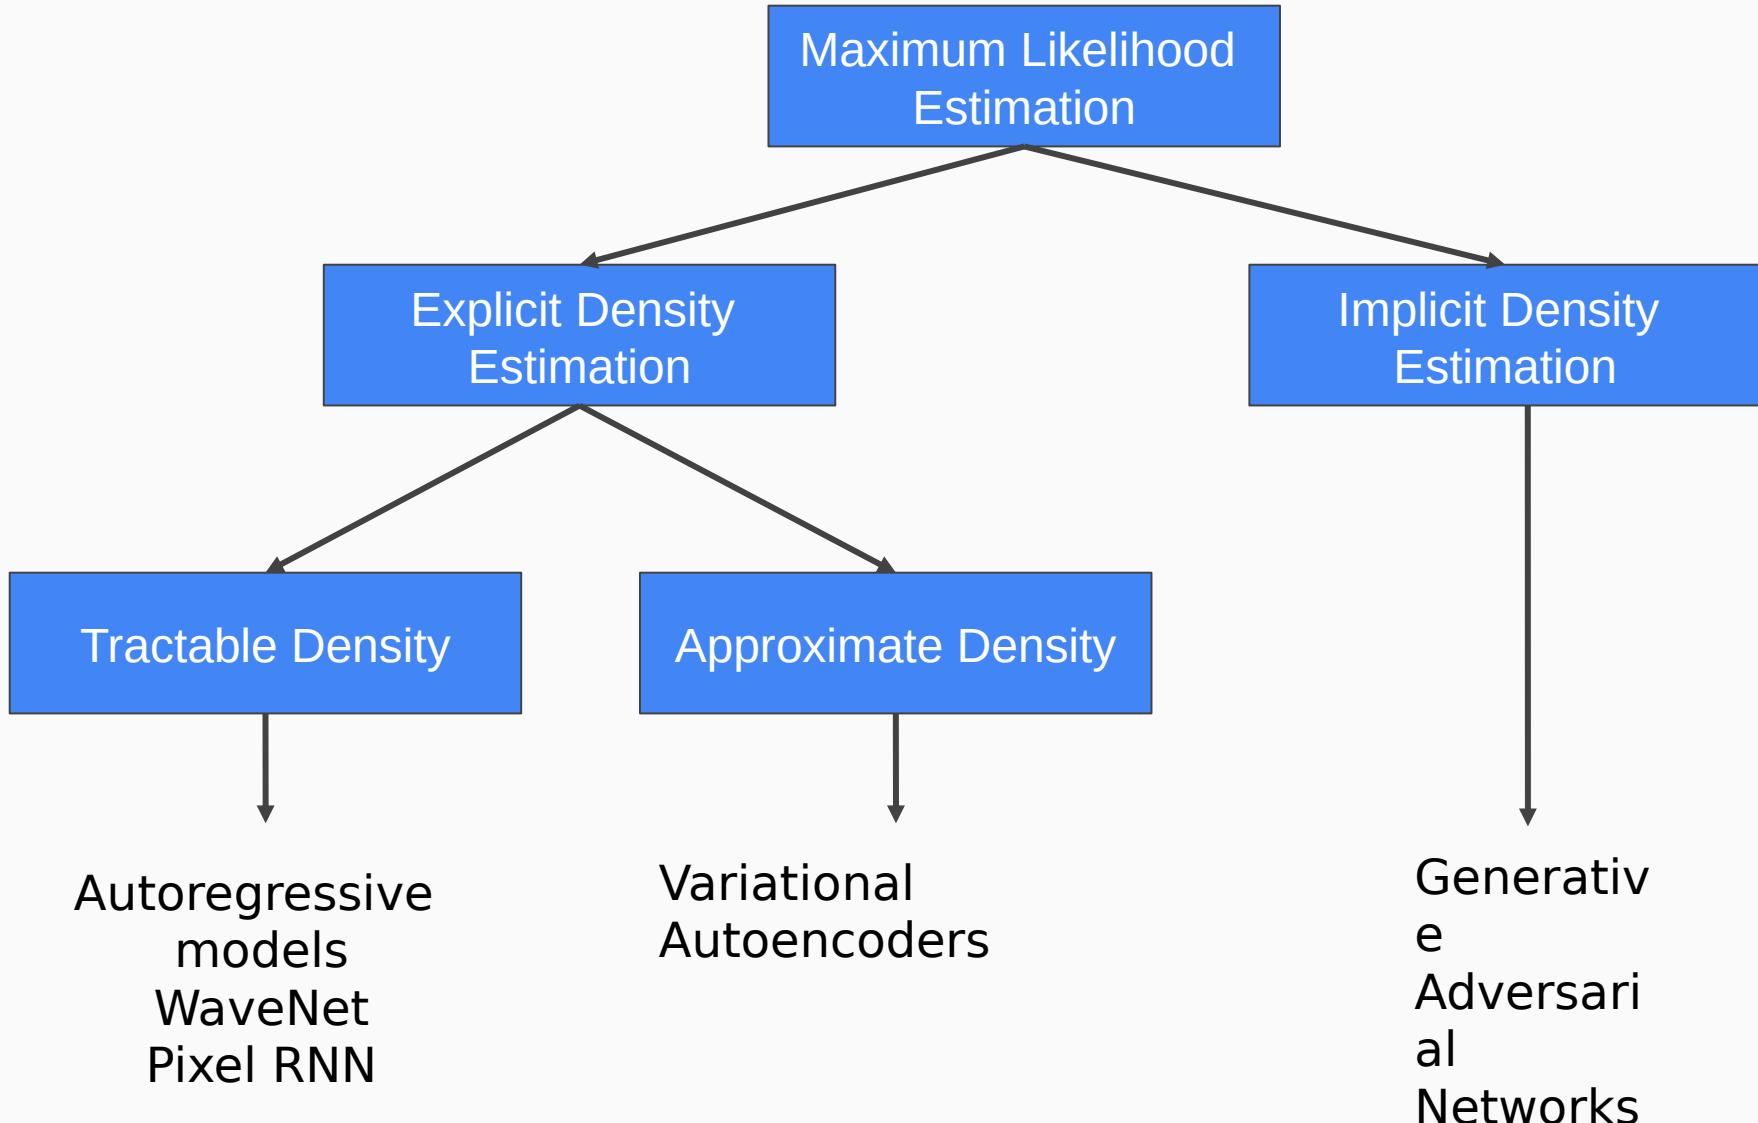
Our goal: Create a model, from which we can sample, that approximates the joint data distribution

I.e. Minimize the KL divergence between the data-generating distribution and the model distribution

$$\min_{p \in \mathcal{P}_{x,z}} D_{\text{KL}}(p_{\text{data}}(\mathbf{x}) \parallel p(\mathbf{x}))$$



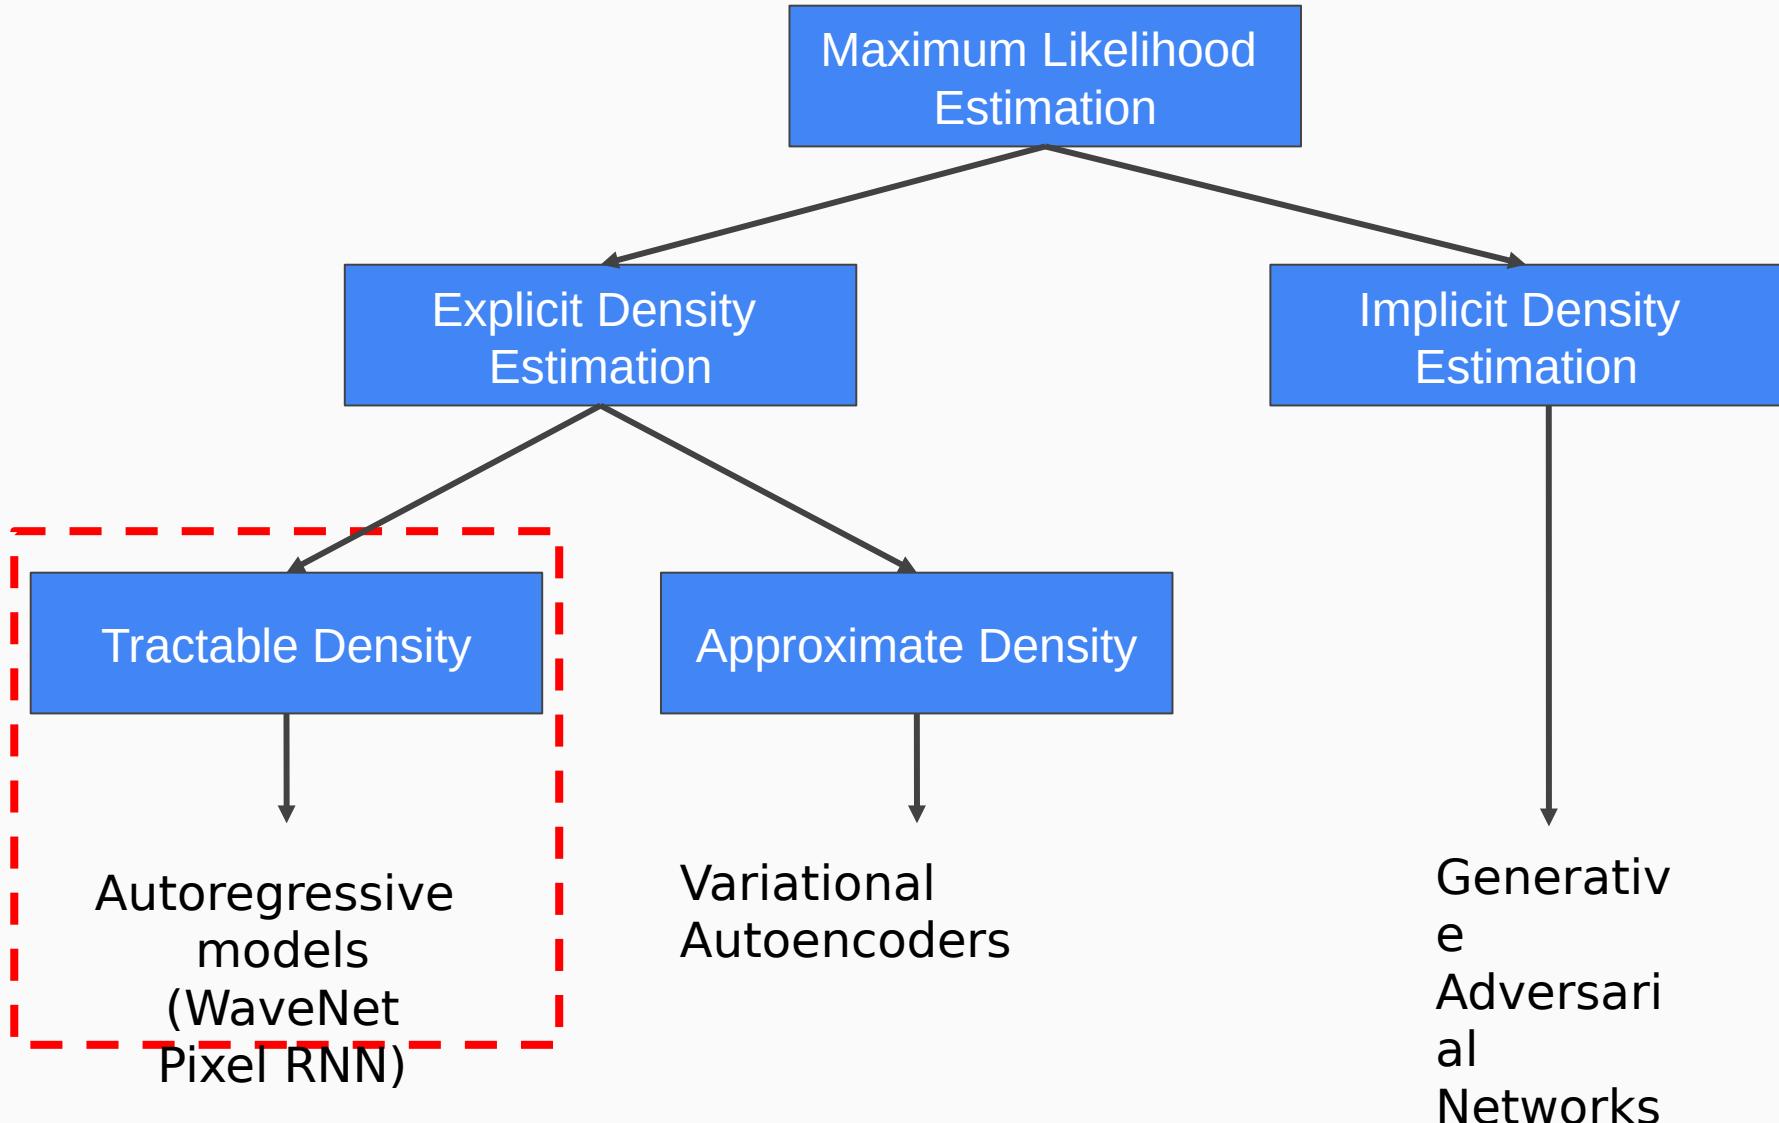
Taxonomy of Generative Models



Adapted from

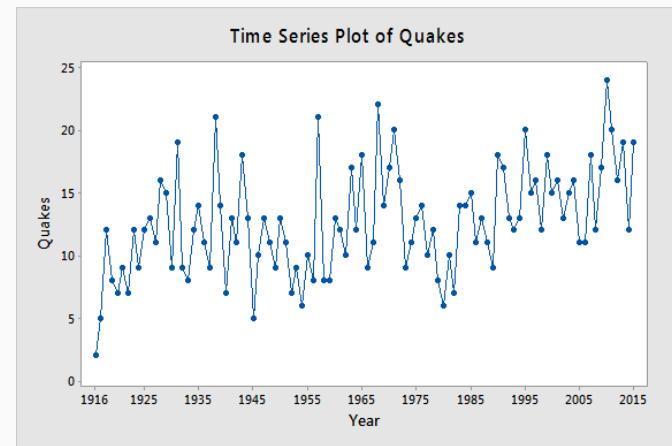
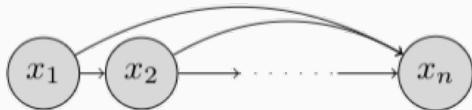
[Ian Goodfellow, Tutorial on Generative Adversarial Networks, NIPS 2016](#)

Taxonomy of Generative Models



Tractable Explicit Density Estimation: Autoregressive Models

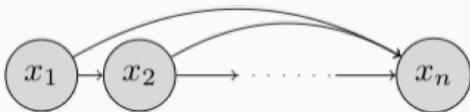
Auto-regression: Prediction based on past data



Source:
[STAT 501, Penn State](#)

Tractable Explicit Density Estimation: Autoregressive Models

Auto-regression: Prediction based on past data

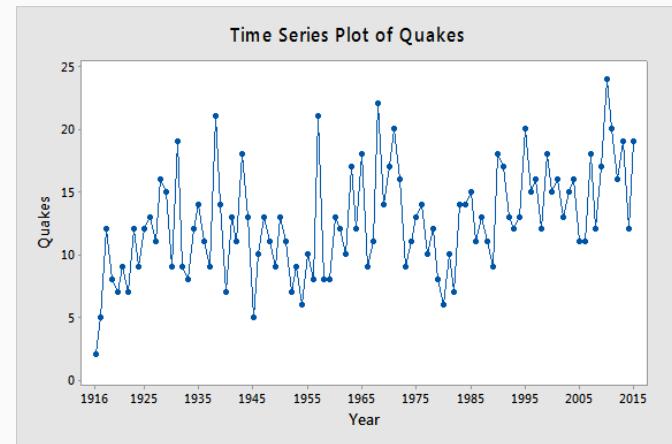


Likelihood of data x
(e.g. image, song)
is defined as a product
of

conditional
probabilities

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^n p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values

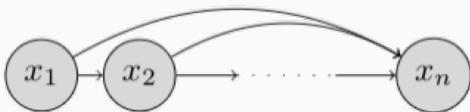


Source:
[STAT 501, Penn State](#)

Training goal: Maximize the likelihood of training data

Tractable Explicit Density Estimation: Autoregressive Models

Auto-regression: Prediction based on past data

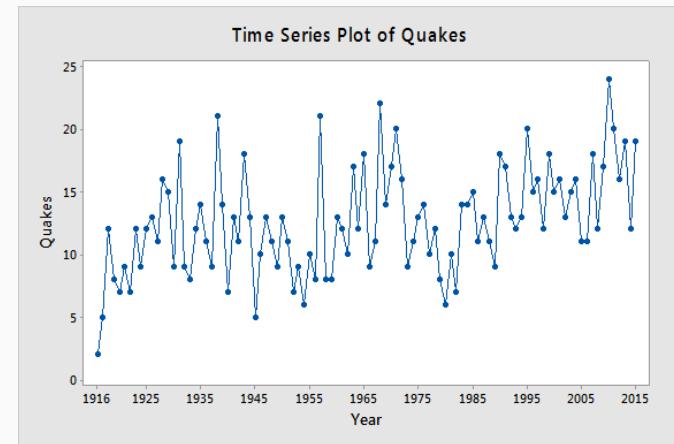


Likelihood of data x
(e.g. image, song)
is defined as a product
of

conditional
probabilities

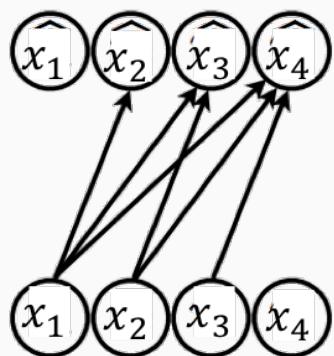
$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^n p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values



Source:
[STAT 501, Penn State](#)

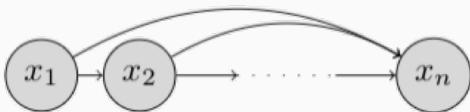
Training goal: Maximize the likelihood of training data



Fully Visible Belief Network
(Source: [Stanford CS236](#))

Tractable Explicit Density Estimation: Autoregressive Models

Auto-regression: Prediction based on past data

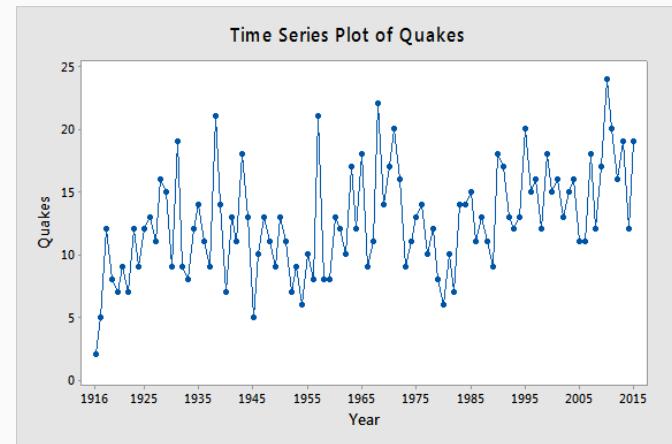


Likelihood of data x (e.g. image, song) is defined as a product of

conditional probabilities

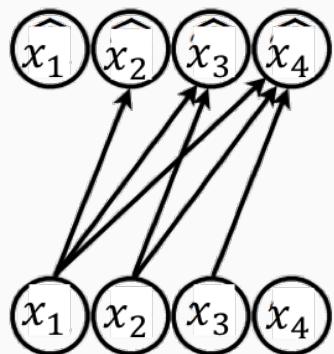
$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^n p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value (e.g. pixel value, audio signal) given all previous values

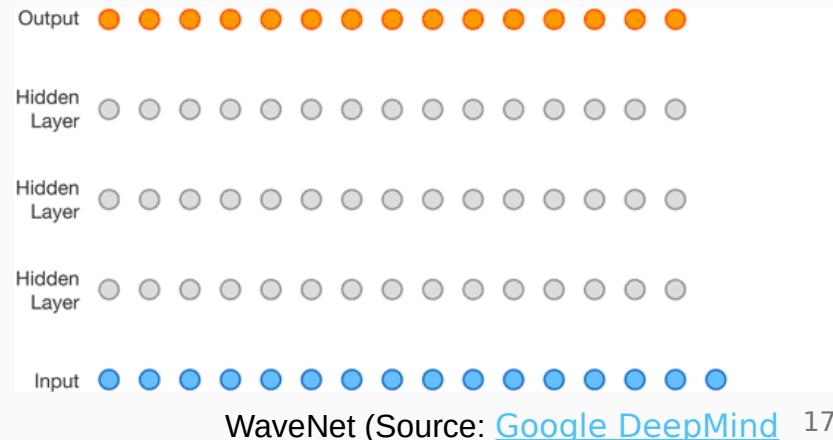


Source:
[STAT 501, Penn State](#)

Training goal: Maximize the likelihood of training data



Fully Visible Belief Network
(Source: [Stanford CS236](#))



WaveNet (Source: [Google DeepMind](#))

Autoregressive Models: PixelRNN

Likelihood of data x
(e.g. image, song)
is defined as a product
of
conditional
probabilitiesⁿ

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^{n^2} p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values

Training goal: Maximize the likelihood of training data

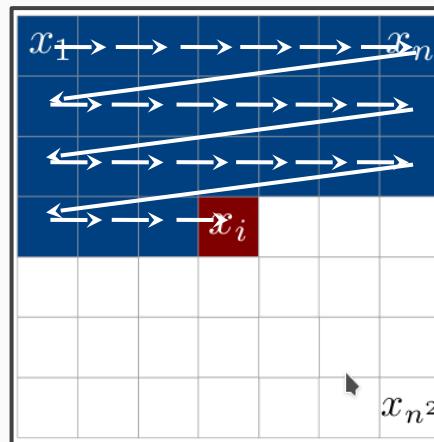
Treat images as a sequence of pixels

Use Recurrent Neural Networks (RNN) model the distribution over pixel values

- LSTM, specifically

Generation begins at a single pixel (e.g. top left) and proceeds sequentially

- Either row by row



Source:
[Oord et al \(2016\), Pixel Recurrent Neural Networks](#)

Autoregressive Models: PixelRNN

Likelihood of data x
(e.g. image, song)
is defined as a product
of
conditional
probabilitiesⁿ

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^{n^2} p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values

Training goal: Maximize the likelihood of training data

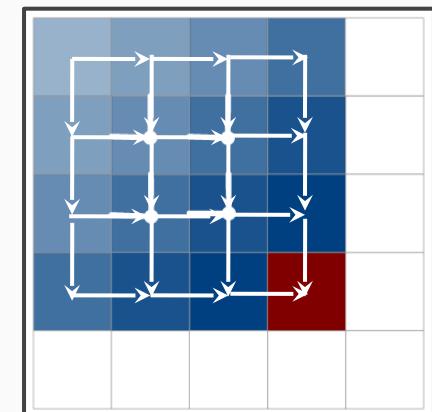
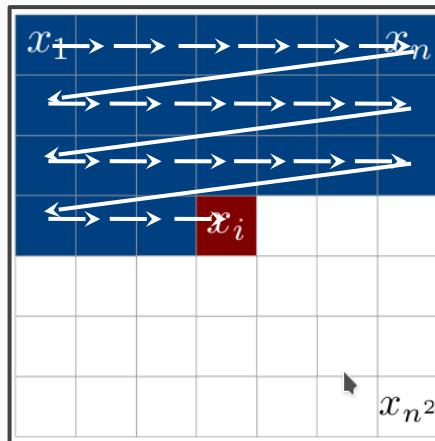
Treat images as a sequence of pixels

Use Recurrent Neural Networks (RNN) model the distribution over pixel values

- LSTM, specifically

Generation begins at a single pixel (e.g. top left) and proceeds sequentially

- Either row by row
- Or bi-directionally



Source:
[Oord et al \(2016\), Pixel Recurrent Neural Networks](#)

Autoregressive Models: PixelRNN

Likelihood of data x
(e.g. image, song)
is defined as a product
of
conditional
probabilitiesⁿ

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^{n^2} p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values

Training goal: Maximize the likelihood of training data

Treat images as a sequence of pixels

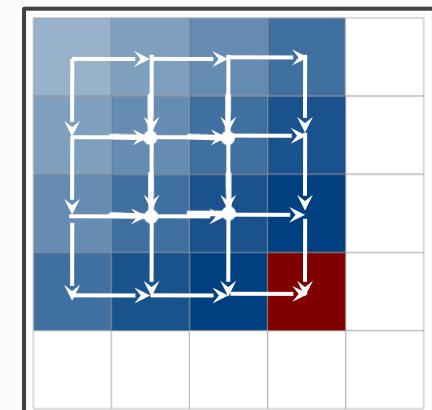
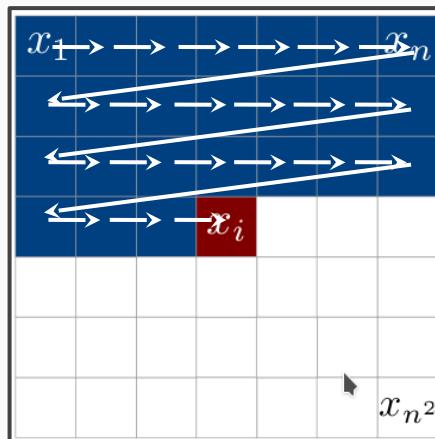
Use Recurrent Neural Networks (RNN) model the distribution over pixel values

- LSTM, specifically

Generation begins at a single pixel (e.g. top left) and proceeds sequentially

- Either row by row

PixelRNN uses a Convolutional Neural Network instead of RNN to model distribution, speeds up



Source:

[Oord et al \(2016\), Pixel Recurrent Neural Networks](#)

Autoregressive Models: PixelRNN

Likelihood of data x
(e.g. image, song)
is defined as a product
of
conditional
probabilitiesⁿ

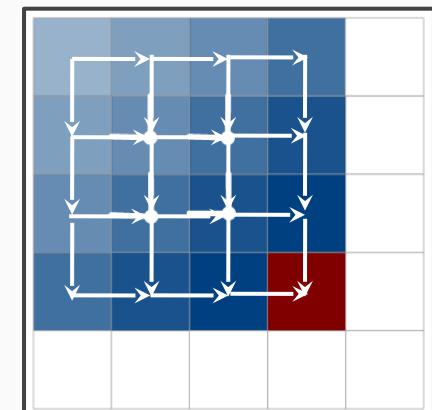
$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1}) = \prod_{i=1}^{n^2} p(x_i | \mathbf{x}_{<i})$$

Probability of i^{th} value
(e.g. pixel value, audio
signal)
given all previous values

Training goal: Maximize the likelihood of training data

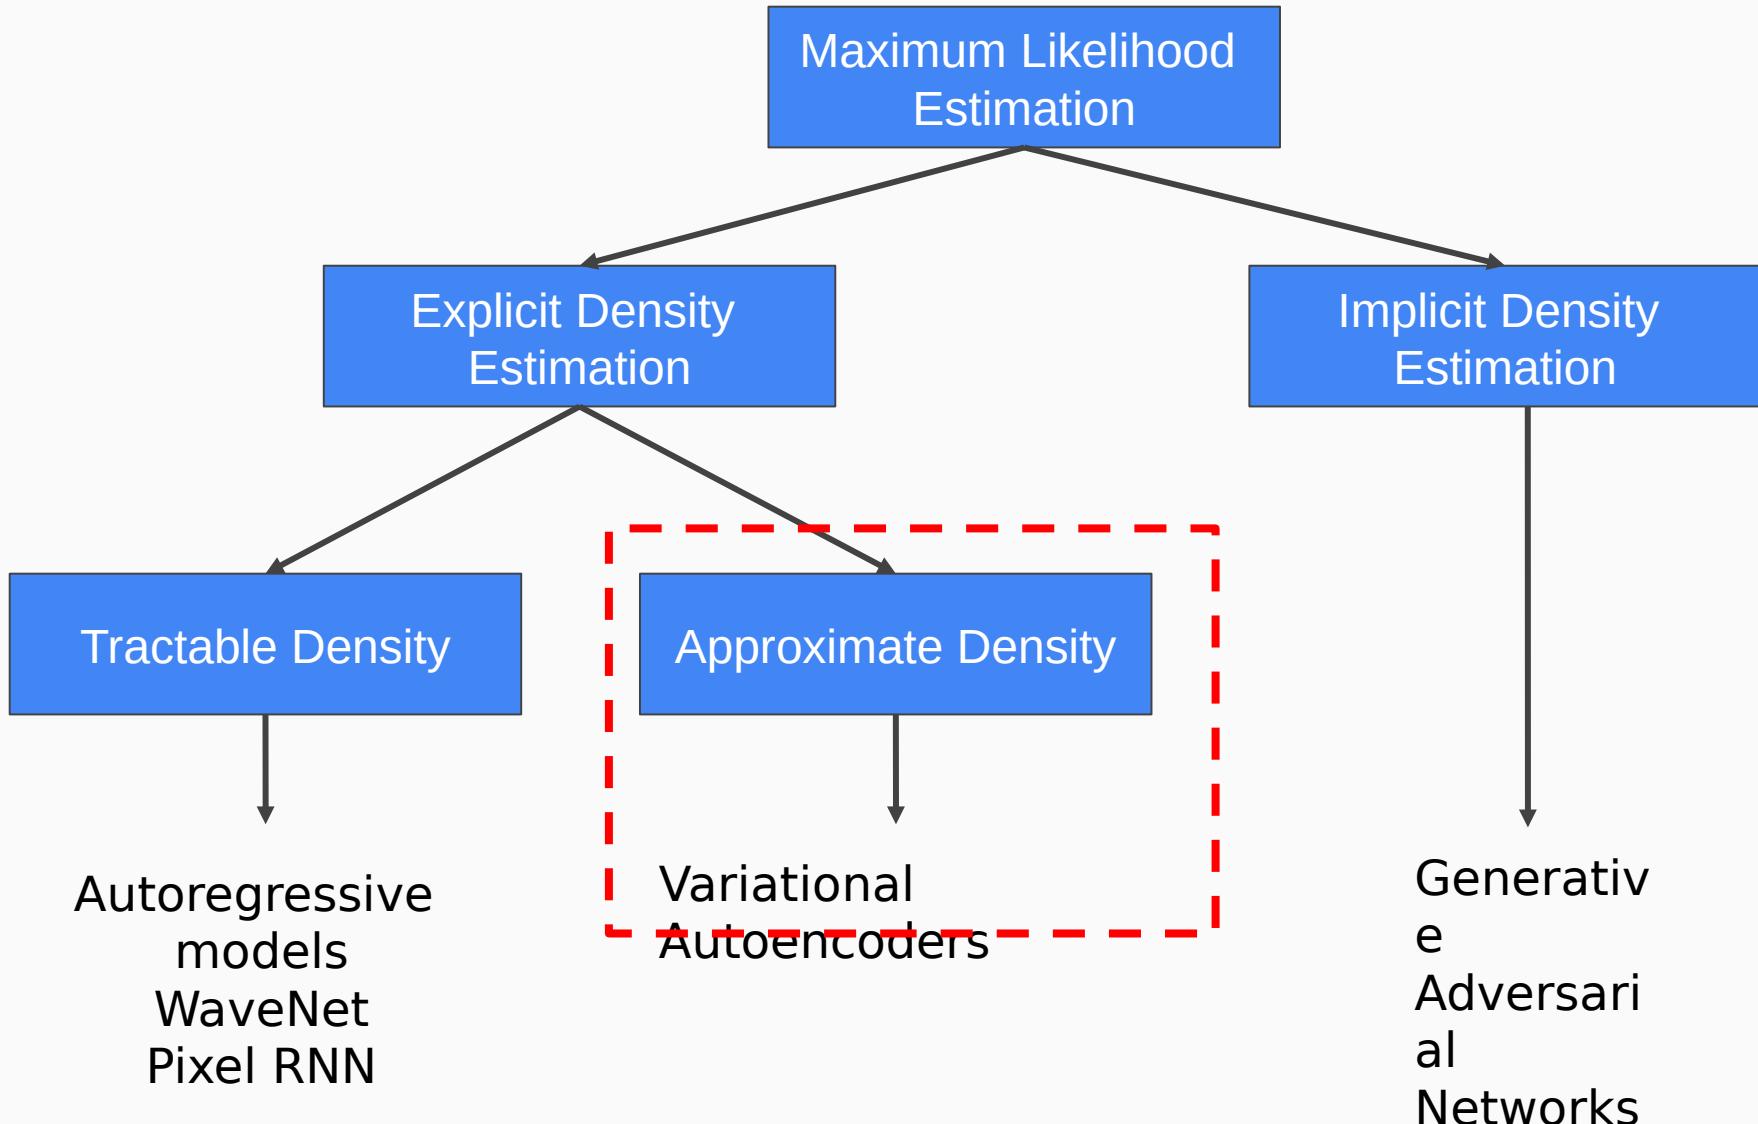


Figure 1. Image completions sampled from a PixelRNN.



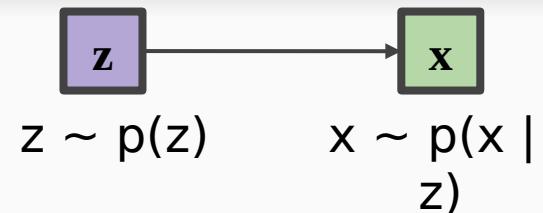
Source:
[Oord et al \(2016\), Pixel Recurrent Neural Networks](#)

Taxonomy of Generative Models



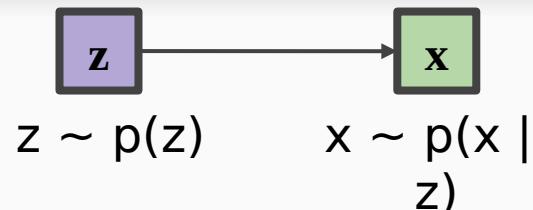
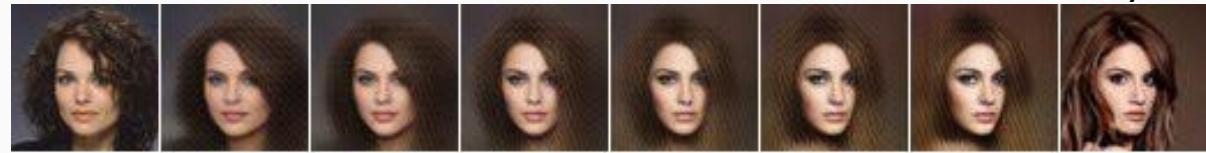
Latent space

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)



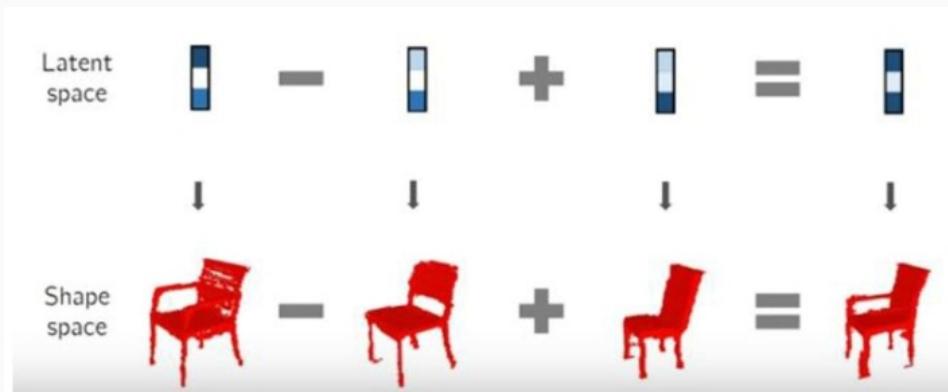
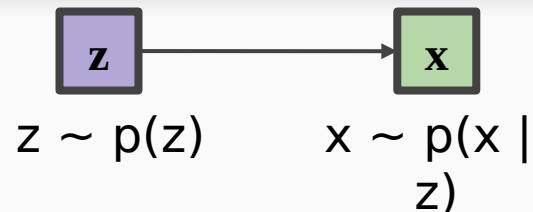
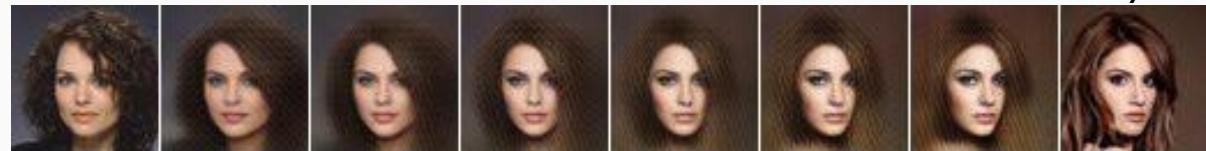
Latent space

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)



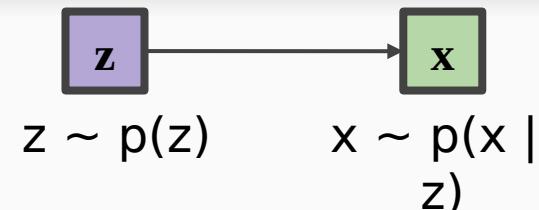
Latent space

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)

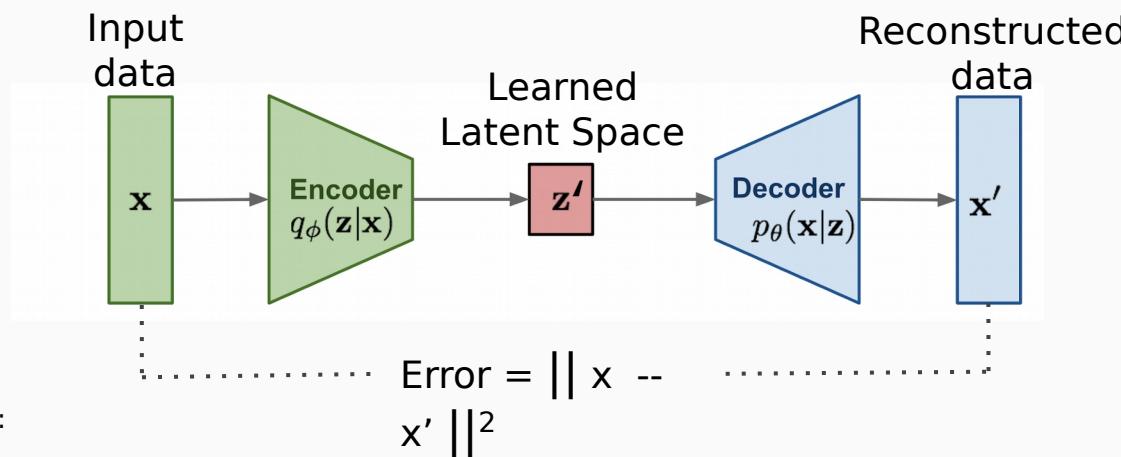


Approximate Explicit Density Estimation: Traditional Autoencoder

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)



Traditional Autoencoder: Unsupervised learning of the latent space of the training data



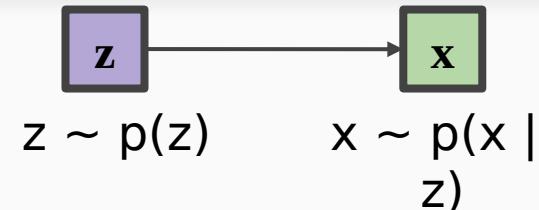
Adapted from:
[Lilian Weng](#)

Training goal: Minimize the loss between the training data and the reconstructed data

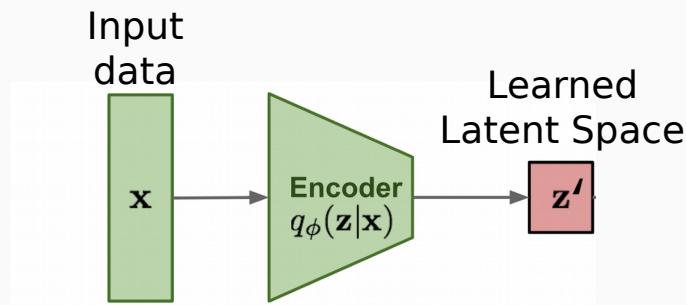
i.e. Train an encoder that creates a latent space which can be used to mimic the

Approximate Explicit Density Estimation: Traditional Autoencoder

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)



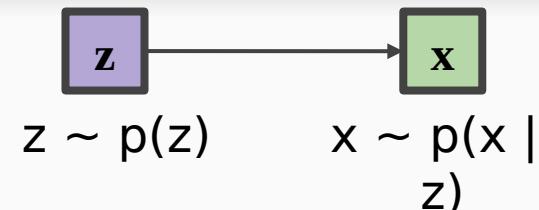
Traditional Autoencoder: Unsupervised learning of the latent space of the training data



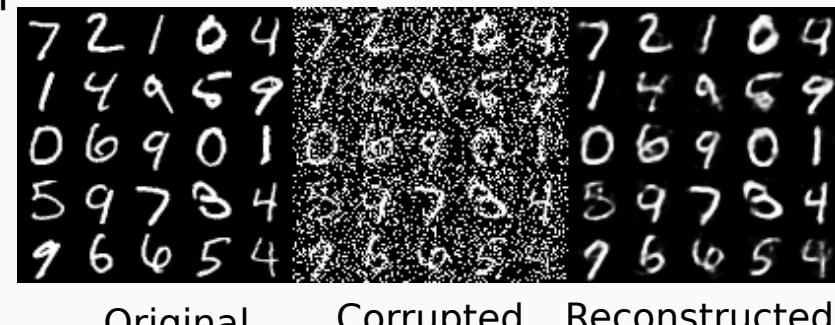
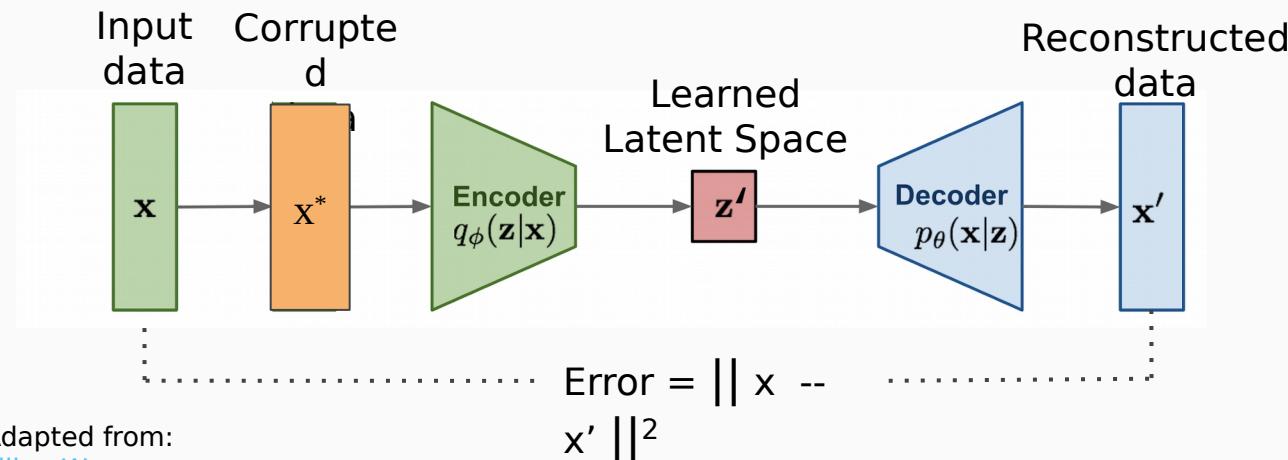
- * Latent space has less dimensions (less features) than the input data
 - Data compression
 - Dimensionality reduction
 - Feature extraction for later supervised learning

Approximate Explicit Density Estimation: Denoising Autoencoder

Latent space: An underlying and unobserved distribution (z) that is assumed to be a generative source of data (x)



Denoising Autoencoder: Reconstruct corrupted images to their original form

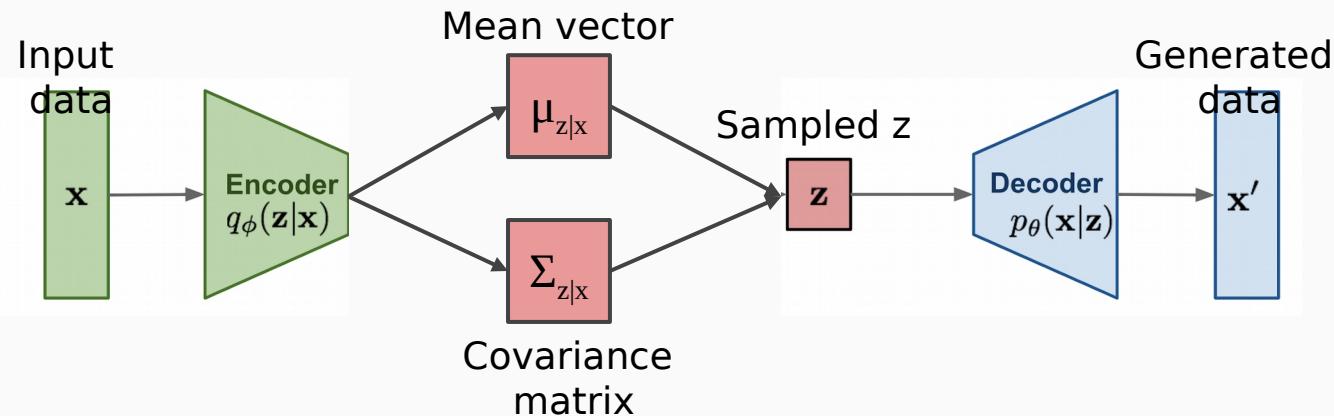


Adapted from:
[Lilian Weng](#)

Approximate Explicit Density Estimation: Variational Autoencoder

Traditional Autoencoder: Map input to **fixed** latent vector z

Variational Autoencoder: Map input to a **distribution** that generates variations in z



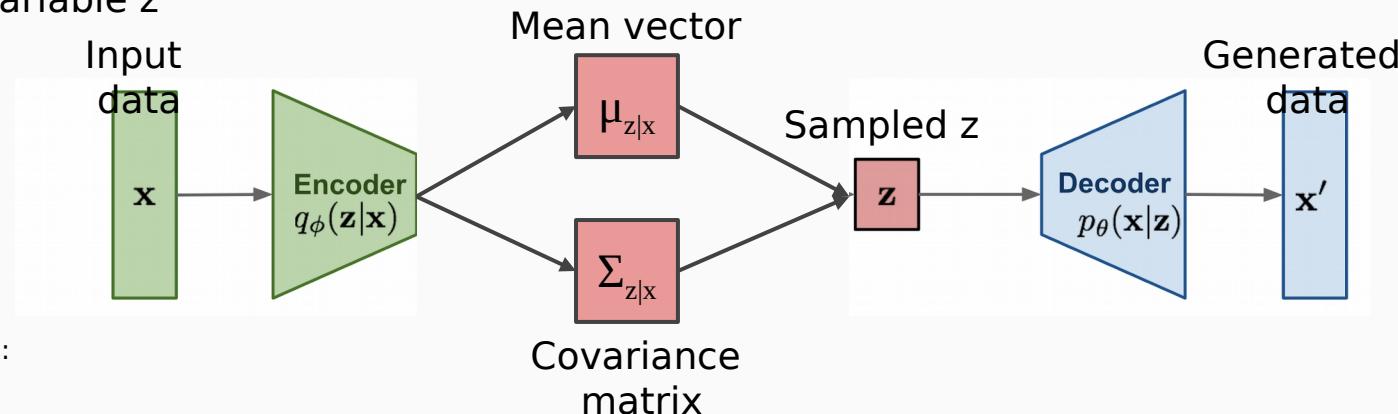
Adapted from:
[Lilian Weng](#)

Approximate Explicit Density Estimation: Variational Autoencoder

Traditional Autoencoder: Map input to **fixed** latent vector z

Variational Autoencoder: Map input to a **distribution** that generates variations in z

- Choose a prior probability for $p(z) \rightarrow$ Often set to a simple distribution i.e. Gaussian
- Need to learn parameters of Decoder Network $p_\theta(x | z)$ that can generate data from a variable z



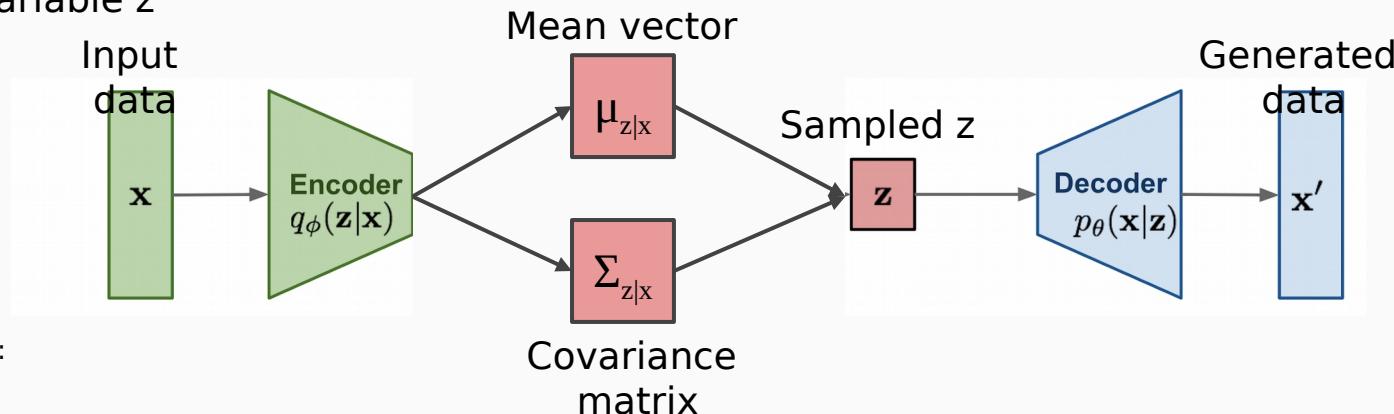
Adapted from:
[Lilian Weng](#)

Approximate Explicit Density Estimation: Variational Autoencoder

Traditional Autoencoder: Map input to **fixed** latent vector z

Variational Autoencoder: Map input to a **distribution** that generates variations in z

- Choose a prior probability for $p(z) \rightarrow$ Often set to a simple distribution i.e. Gaussian
- Need to learn parameters of Decoder Network $p_\theta(x | z)$ that can generate data from a variable z



Adapted from:
[Lilian Weng](#)

$$\min_{p \in \mathcal{P}_{\mathbf{x}, \mathbf{z}}} D_{\text{KL}}(p_{\text{data}}(\mathbf{x}) \| p(\mathbf{x})) = \max_{p \in \mathcal{P}_{\mathbf{x}, \mathbf{z}}} \sum_{\mathbf{x} \in \mathcal{D}} \log p(\mathbf{x})$$

Data
likelihood:

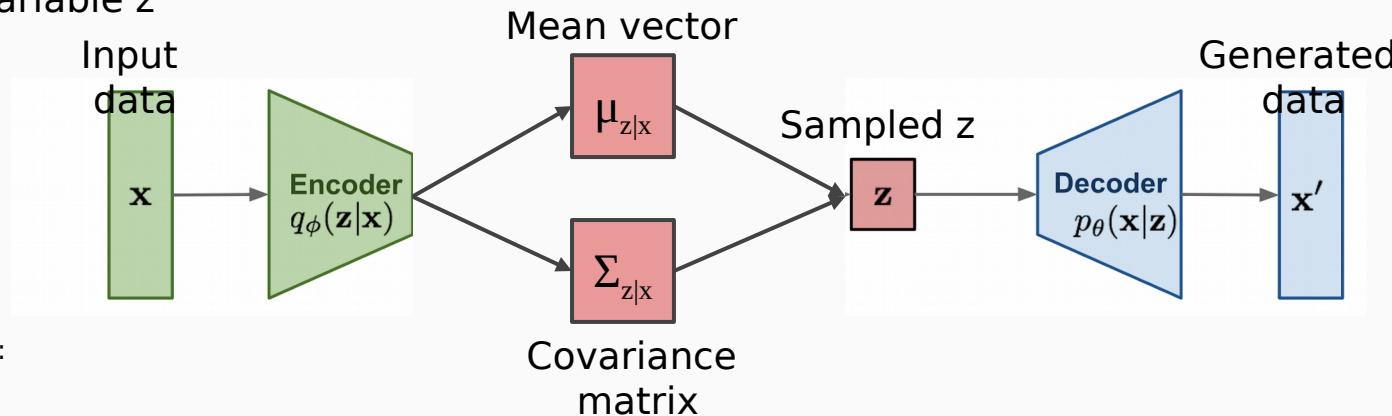
$$p_\theta(x) = \int p_\theta(z)p_\theta(x|z)dz$$

Approximate Explicit Density Estimation: Variational Autoencoder

Traditional Autoencoder: Map input to **fixed** latent vector z

Variational Autoencoder: Map input to a **distribution** that generates variations in z

- Choose a prior probability for $p(z) \rightarrow$ Often set to a simple distribution i.e. Gaussian
- Need to learn parameters of Decoder Network $p_\theta(x | z)$ that can generate data from a variable z



Adapted from:
[Lilian Weng](#)

$$\min_{p \in \mathcal{P}_{\mathbf{x}, \mathbf{z}}} D_{\text{KL}}(p_{\text{data}}(\mathbf{x}) \| p(\mathbf{x})) = \max_{p \in \mathcal{P}_{\mathbf{x}, \mathbf{z}}} \sum_{\mathbf{x} \in \mathcal{D}} \log p(\mathbf{x})$$

Data

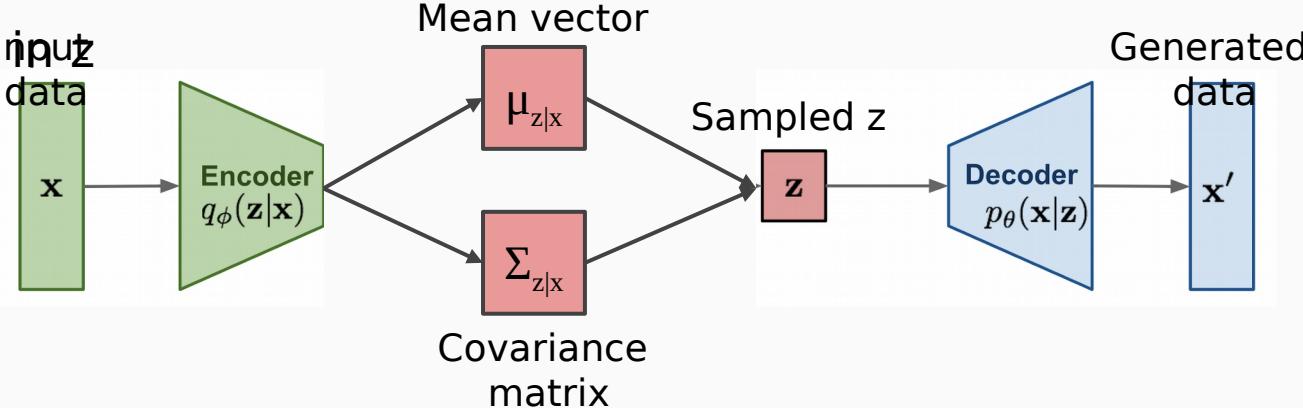
likelihood:

$$p_\theta(x) = \int p_\theta(z)p_\theta(x|z)dz$$

Intractable (computationally expensive) for high dimensional z (which is the case for most problems of interest)

Approximate Explicit Density Estimation: Variational Autoencoder

Variational Autoencoder: Map input to a **distribution** that generates variations



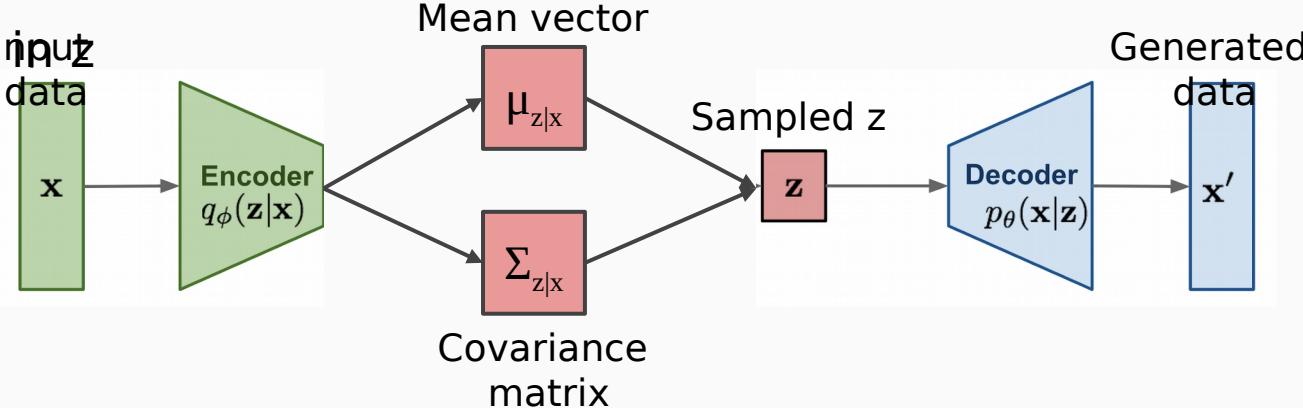
Adapted from:
[Lilian Weng](#)

$$\log p_\theta(x^{(i)}) = \mathbb{E}_z \left[\log p_\theta(x^{(i)} | z) \right] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z)) + D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))$$

Solvable **Intractable, but know it
is ≥ 0**

Approximate Explicit Density Estimation: Variational Autoencoder

Variational Autoencoder: Map input to a **distribution** that generates variations



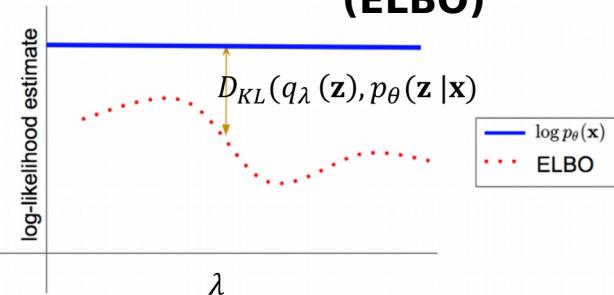
Adapted from:
[Lilian Weng](#)

$$\log p_\theta(x^{(i)}) = \mathbb{E}_z \left[\log p_\theta(x^{(i)} | z) \right] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z)) + D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))$$

Solvable **Intractable, but know it is ≥ 0**

Variational Lower Bound (ELBO)

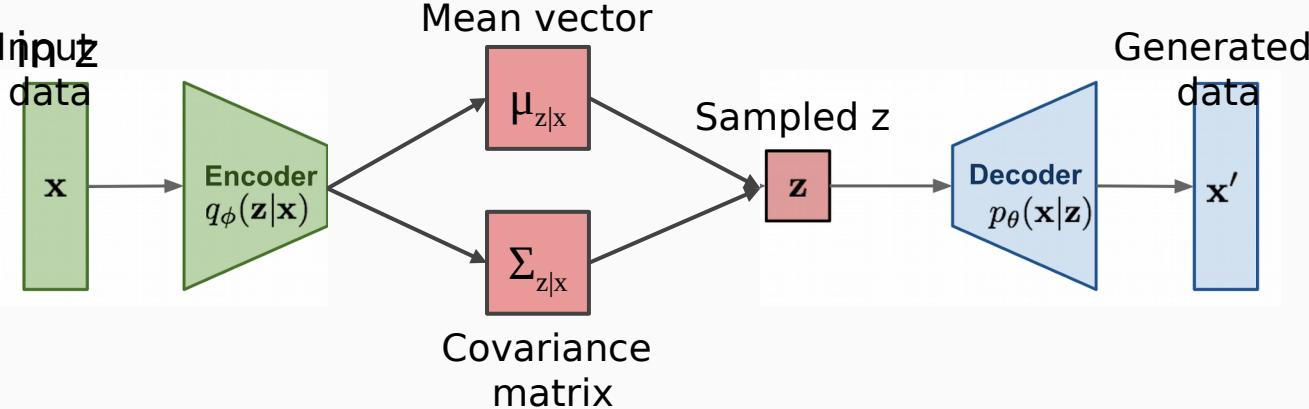
$$= \mathbb{E}_z \left[\log p_\theta(x^{(i)} | z) \right] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))$$



Training goal: Maximize the Variational Lower Bound

Approximate Explicit Density Estimation: Variational Autoencoder

Variational Autoencoder: Map input to a **distribution** that generates variations



Adapted from:
[Lilian Weng](#)

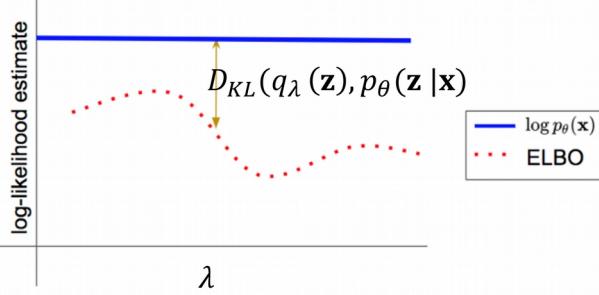
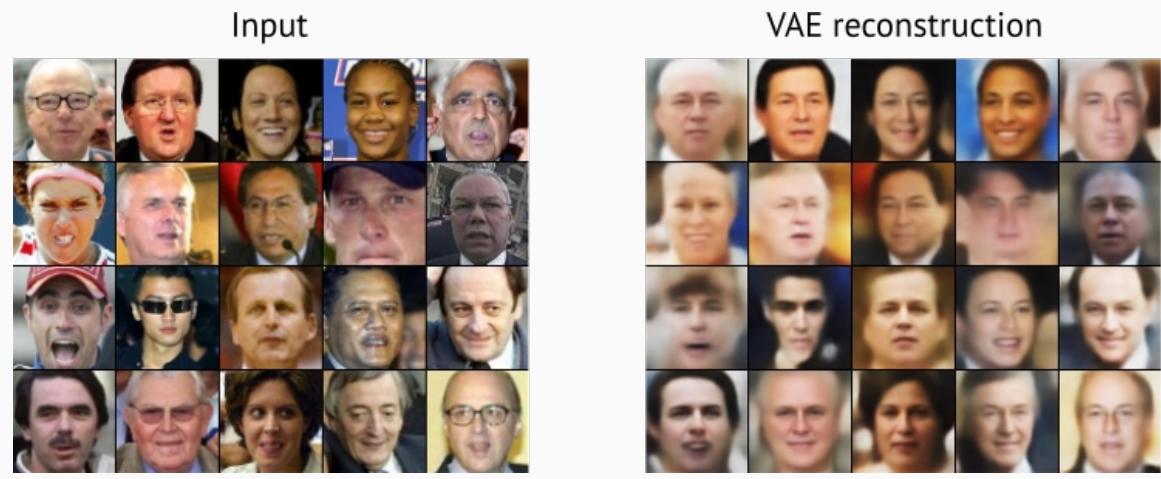
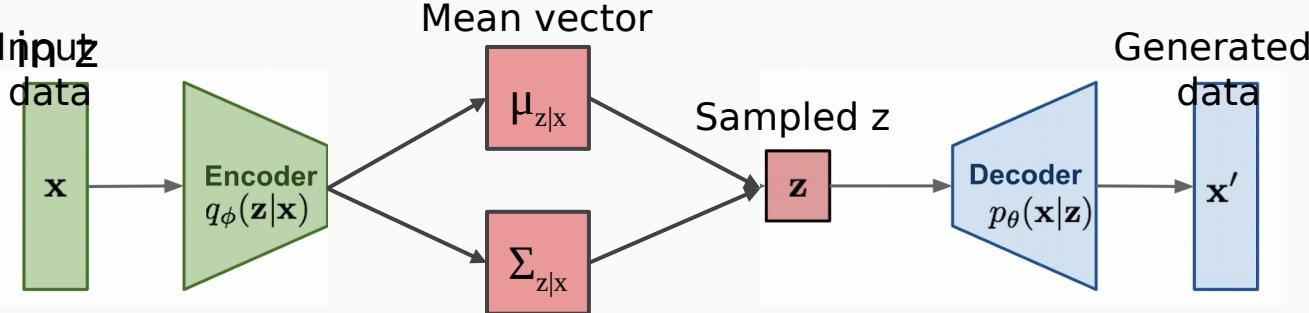


Image Source

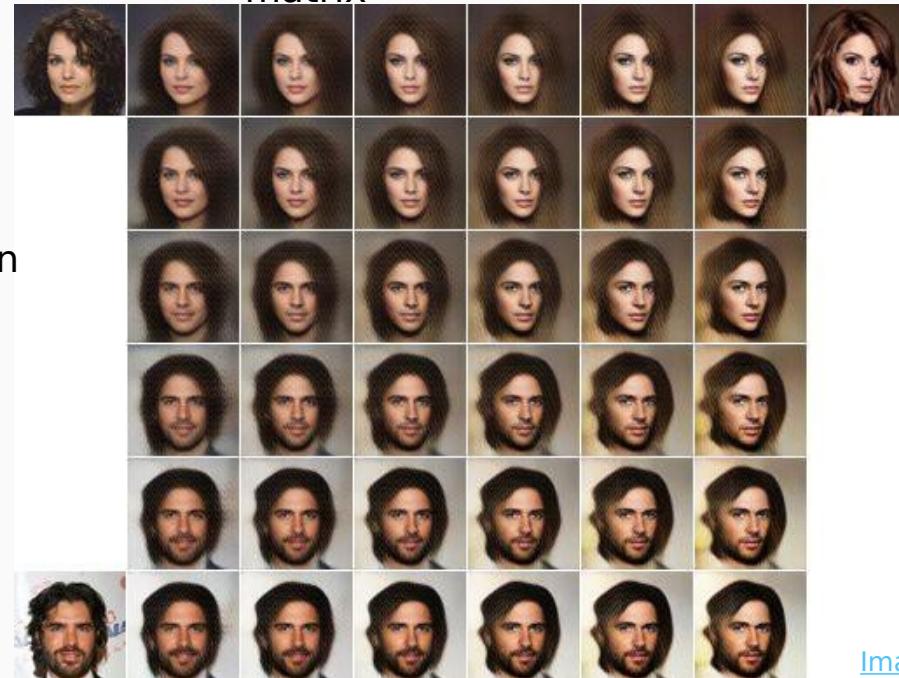
Approximate Explicit Density Estimation: Variational Autoencoder

Variational Autoencoder: Map input to a **distribution** that generates variations



Adapted from:
[Lilian Weng](#)

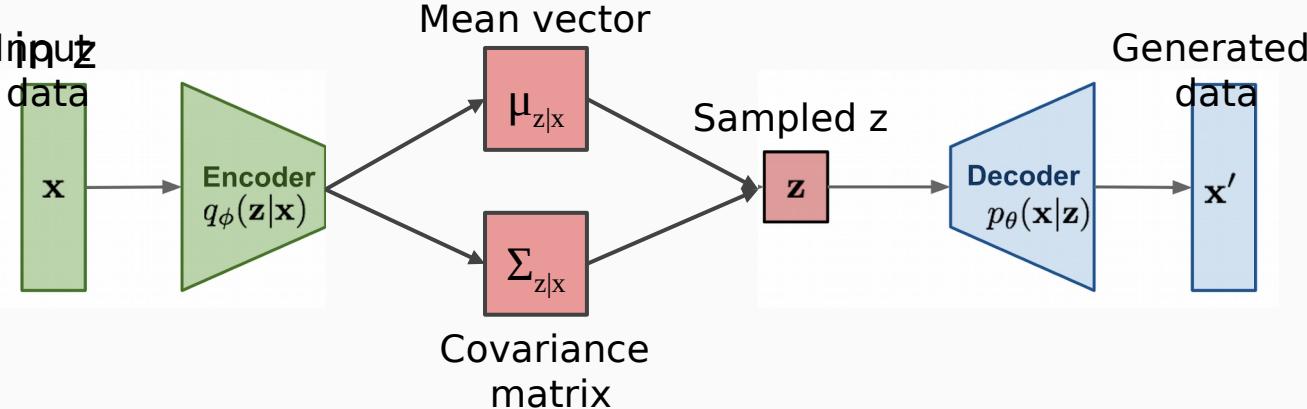
Smooth interpolations in
latent space



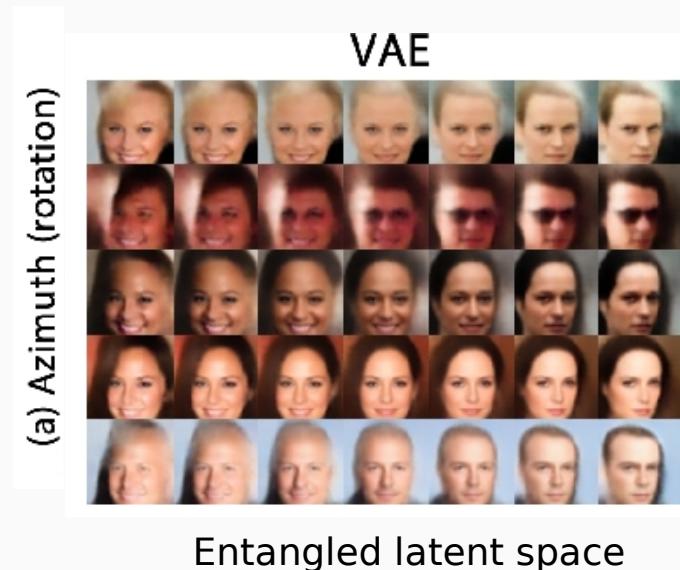
[Image Source](#)

Approximate Explicit Density Estimation: Variational Autoencoder

Variational Autoencoder: Map input to a **distribution** that generates variations

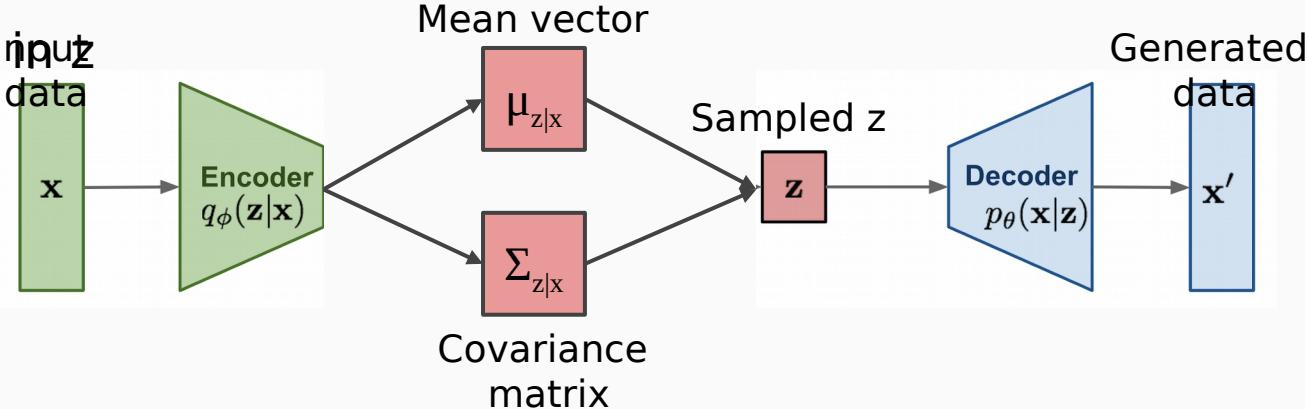


Adapted from:
[Lilian Weng](#)

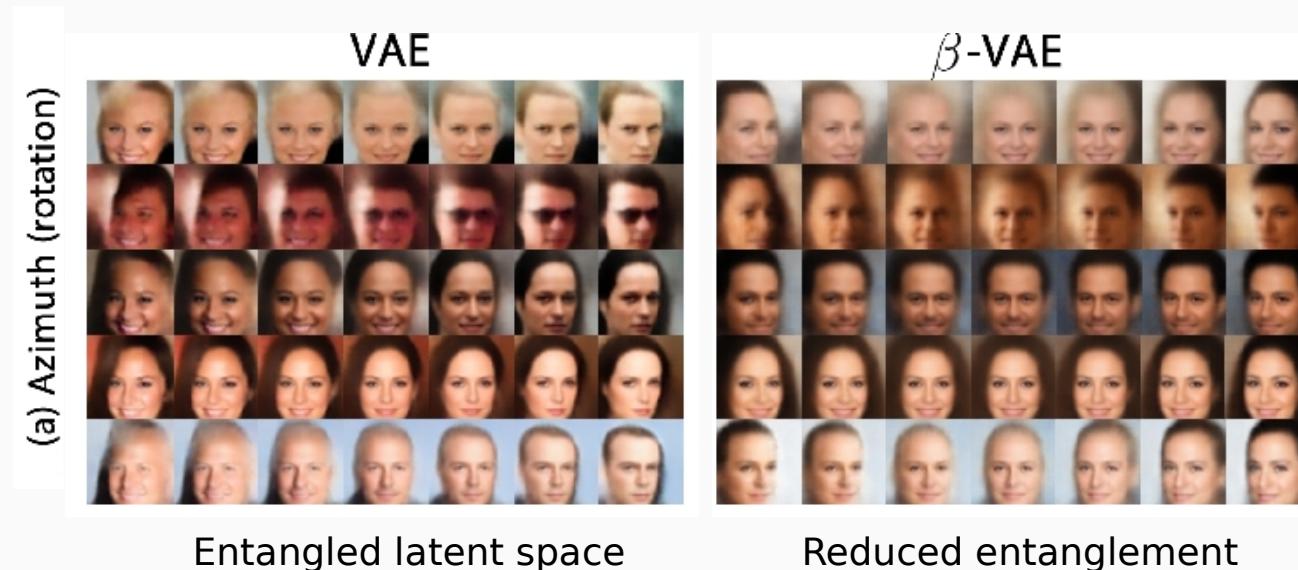


Approximate Explicit Density Estimation: Variational Autoencoder

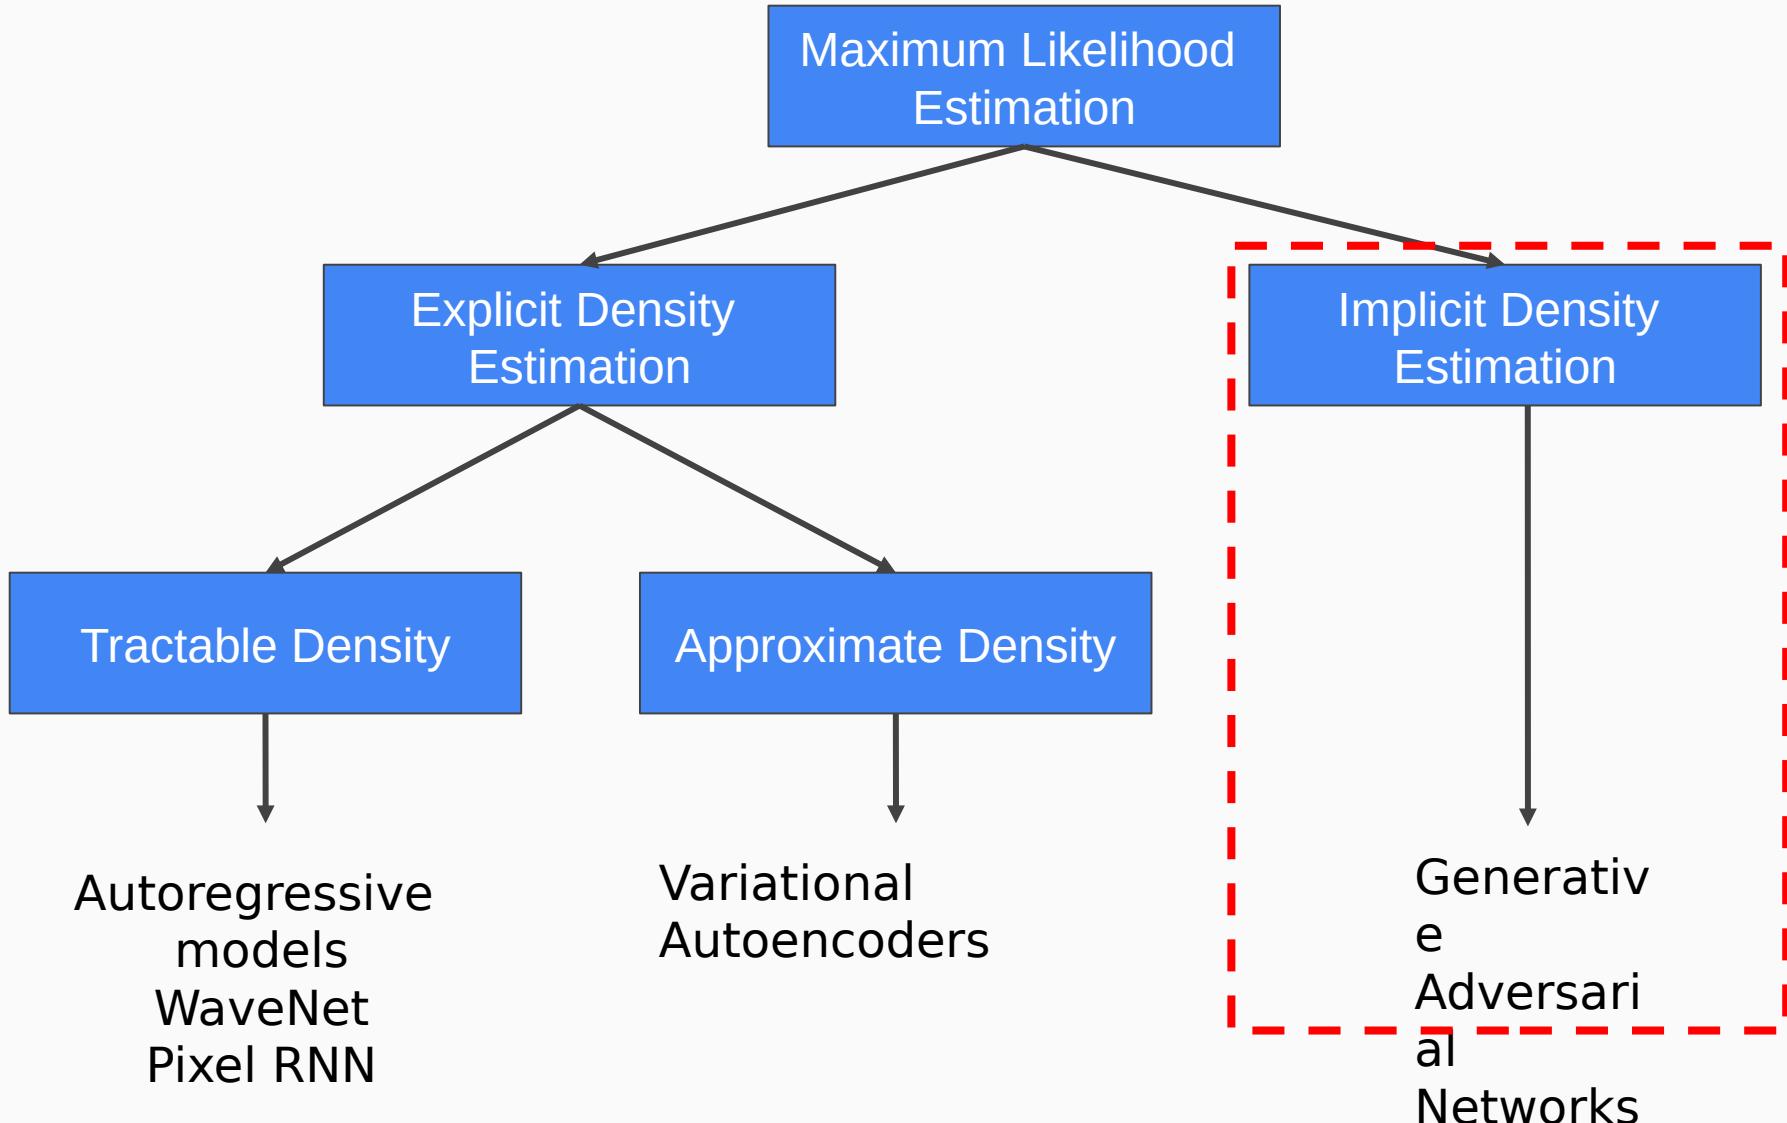
Variational Autoencoder: Map input to a **distribution** that generates variations



Adapted from:
[Lilian Weng](#)

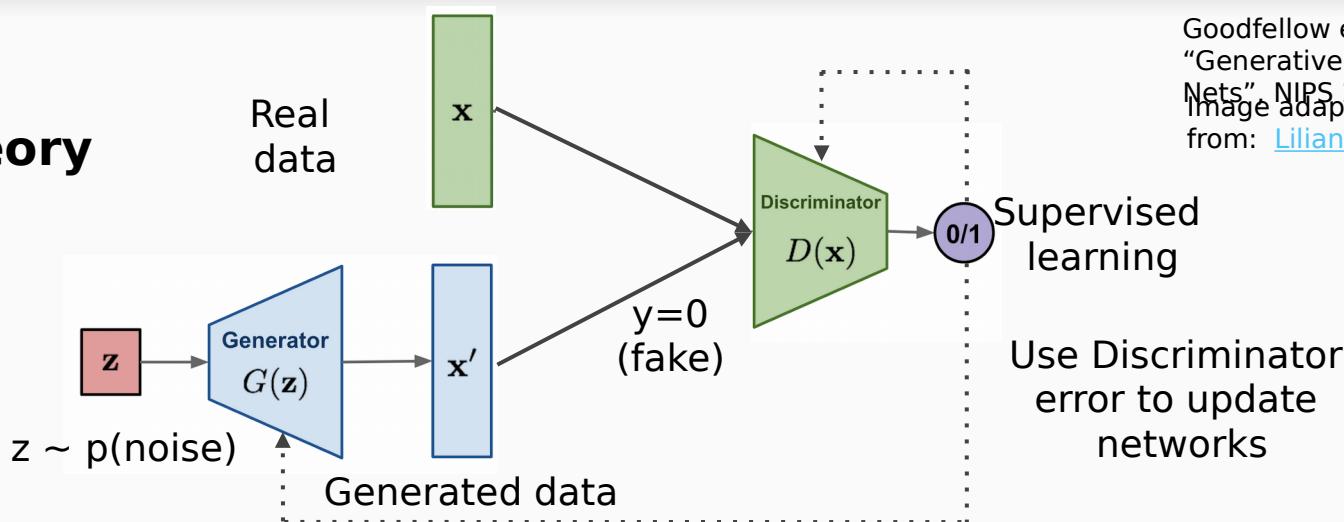


Taxonomy of Generative Models



Implicit Density Estimation: Generative Adversarial Networks

Game theory



Goodfellow et al,
“Generative Adversarial
Nets”, NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

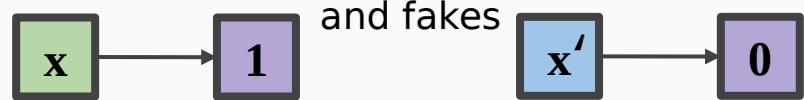
Learns to transform noise into fakes in order to fool the Discriminator



vs

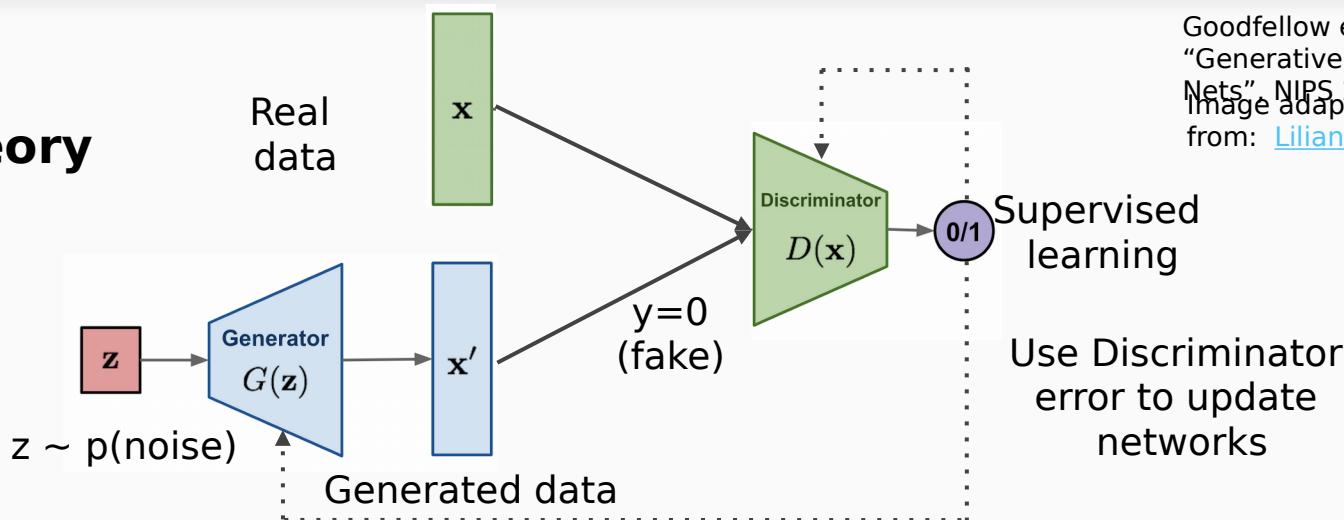
Discriminator

Tries to tell the difference between real examples (data) and fakes



Implicit Density Estimation: Generative Adversarial Networks

Game theory



Goodfellow et al,
“Generative Adversarial
Nets”, NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

Learns to transform noise into fakes in order to fool the Discriminator



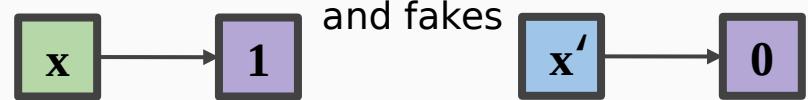
$$\min_{\theta} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Generator tries to minimize the Discriminator's accuracy on generated data

vs

Discriminator

Tries to tell the difference between real examples (data) and fakes

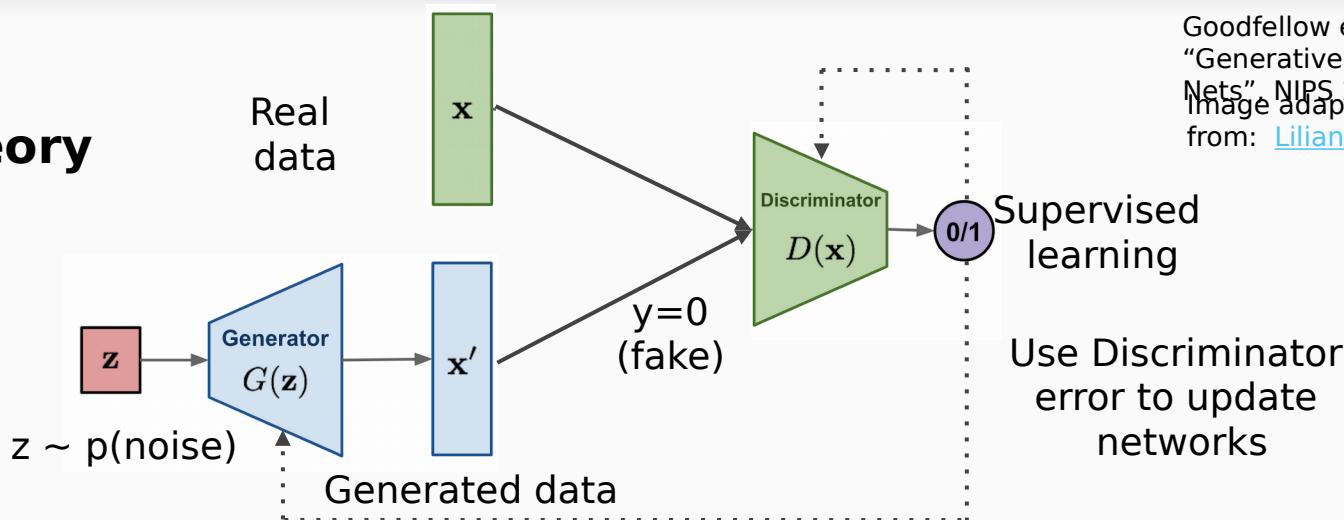


$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Discriminator tries to maximize its accuracy on identifying real as real and fake as fake

Implicit Density Estimation: Generative Adversarial Networks

Game theory



Goodfellow et al,
“Generative Adversarial
Nets”, NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

Learns to transform noise into fakes in order to fool the Discriminator



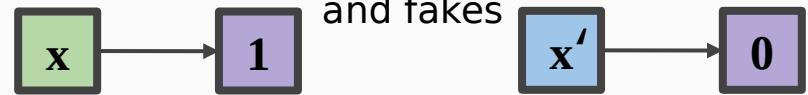
$$\min_{\theta} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Generator tries to minimize the Discriminator's accuracy on generated data

vs

Discriminator

Tries to tell the difference between real examples (data) and fakes



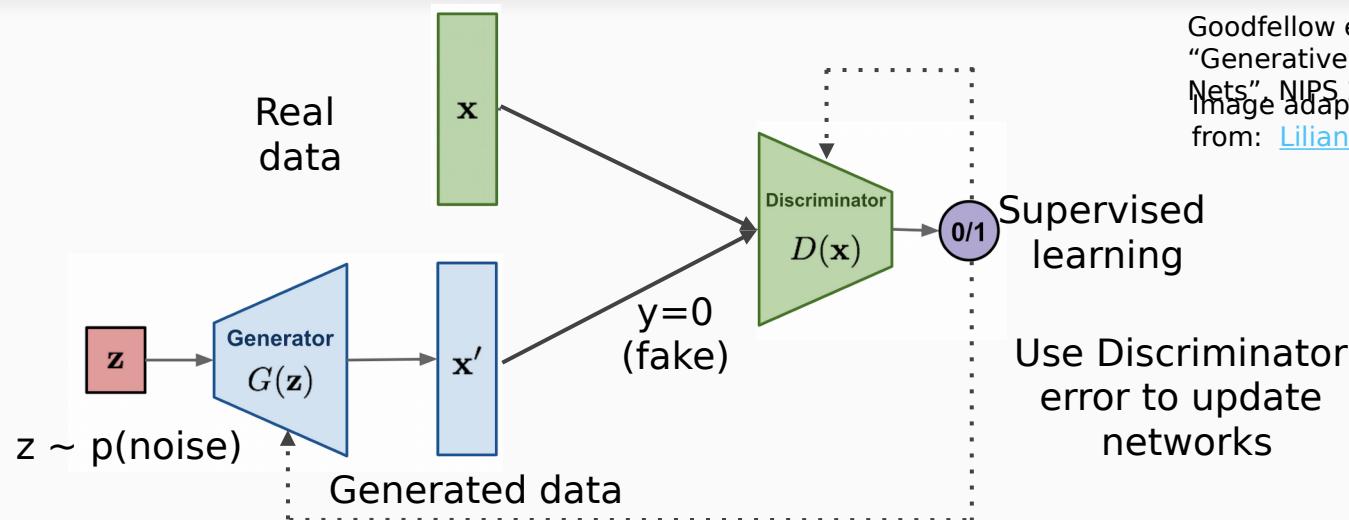
$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Discriminator tries to maximize its accuracy on identifying real as real and fake as fake

AN objective function

$$\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

Implicit Density Estimation: Generative Adversarial Networks



Goodfellow et al,
"Generative Adversarial
Nets", NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

Learns to transform noise into fakes in order to fool the Discriminator



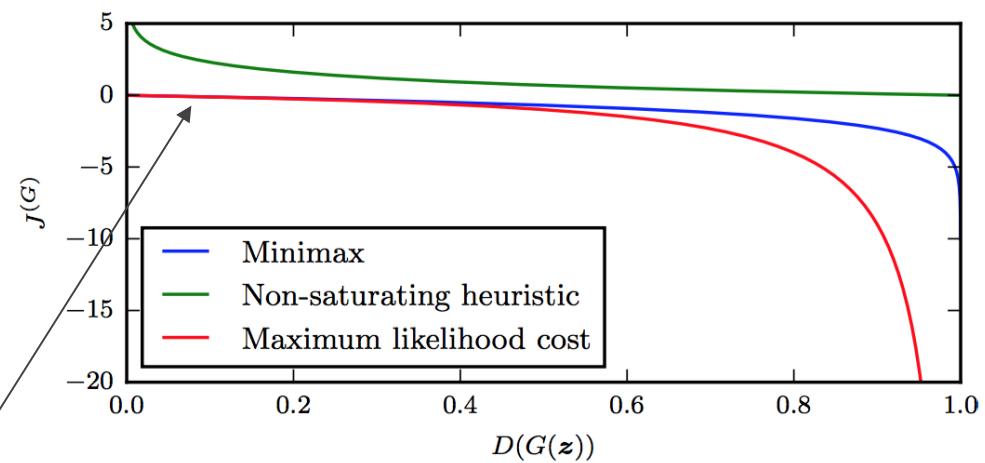
$$\min_{\theta} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Generator tries to minimize the Discriminator's accuracy on generated data

* Difficult to train in this set-up because the gradients for Generator are very close to 0 at the start **

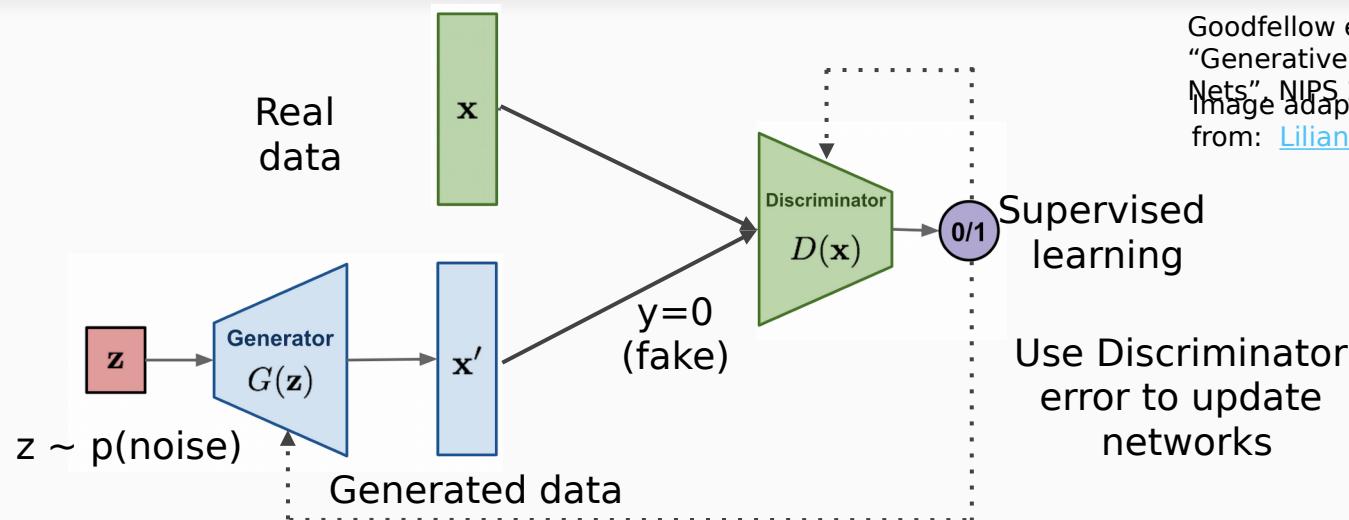
AN objective function

$$\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$



Take as take

Implicit Density Estimation: Generative Adversarial Networks



Goodfellow et al,
"Generative Adversarial
Nets", NIPS 2014
Image adapted
from: [Lilian Weng](#)

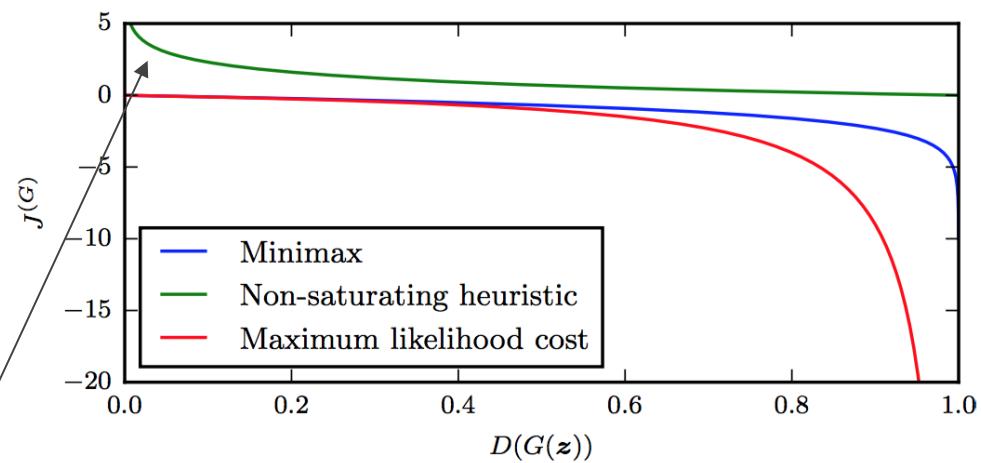
Generator

Learns to transform noise into fakes in order to fool the Discriminator



$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

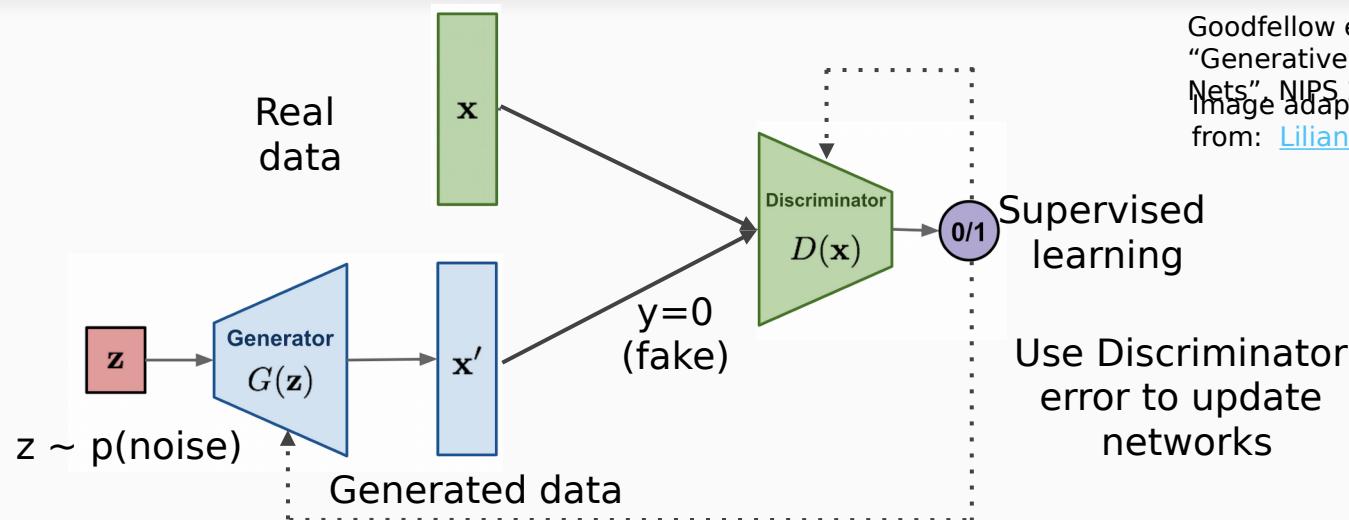
i.e. Generator tries to maximize the Discriminator's error on generated data



**** Higher gradient signals for bad examples, works better in practice ****

AN objective function $\min_{\theta} \max_{\phi} V(G_{\theta}, D_{\phi}) = \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$

Implicit Density Estimation: Generative Adversarial Networks



Goodfellow et al,
"Generative Adversarial
Nets", NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

Learns to transform noise into fakes in order to fool the Discriminator



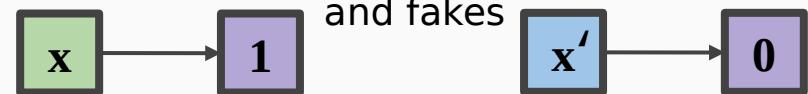
$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

i.e. Generator tries to maximize the Discriminator's error on generated data

vs

Discriminator

Tries to tell the difference between real examples (data) and fakes



$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

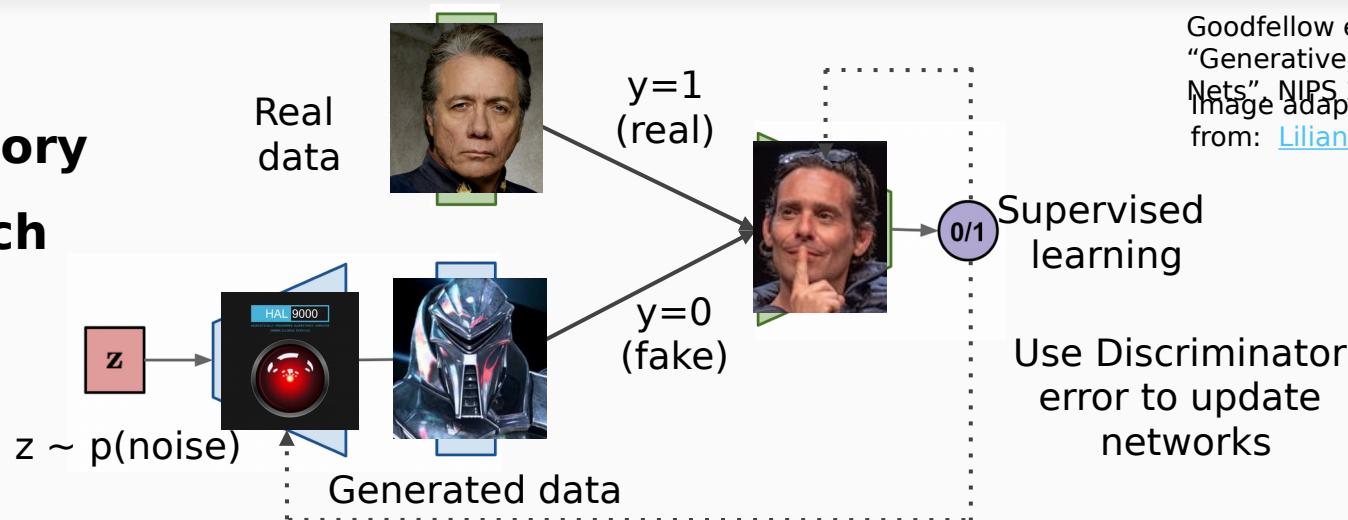
i.e. Discriminator tries to maximize its accuracy on identifying real as real and fake as fake

AN objective function!

Alternate between maximizing Generator and Discriminator objective functions

Implicit Density Estimation: Generative Adversarial Networks

Game theory approach



Generator

Learns to transform noise into fakes in order to fool the Discriminator

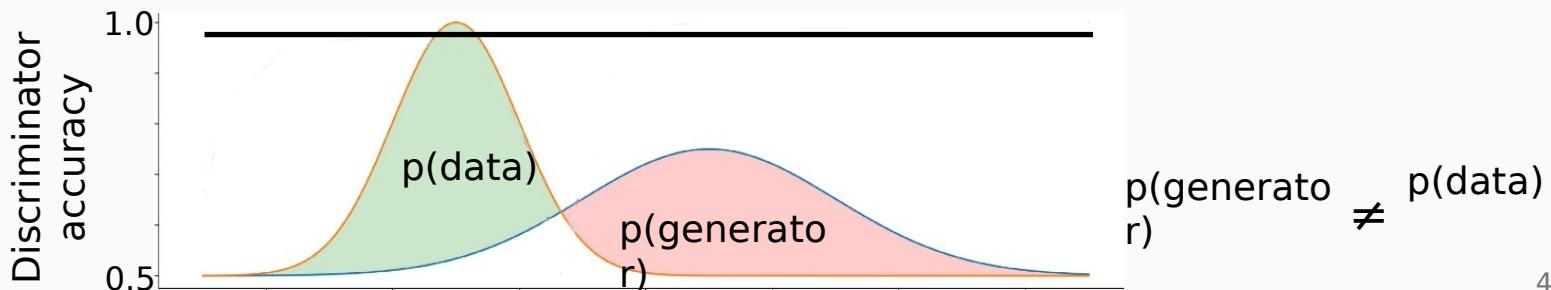
$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

vs

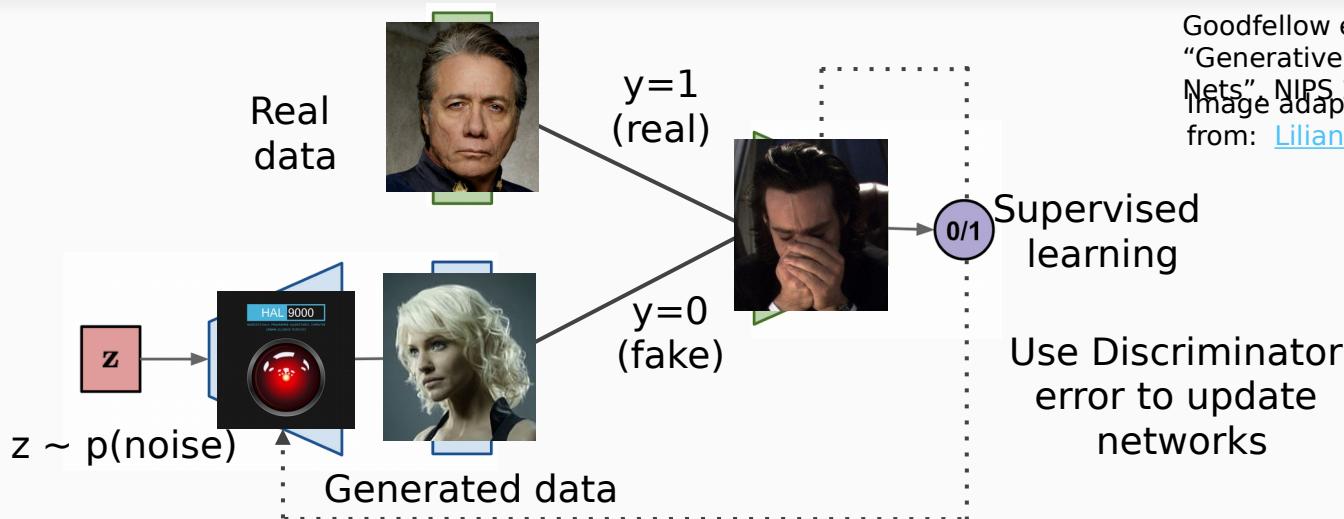
Discriminator

Tries to tell the difference between real examples (data) and fakes

$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$



Implicit Density Estimation: Generative Adversarial Networks



Generator

Learns to transform noise into fakes in order to fool the Discriminator

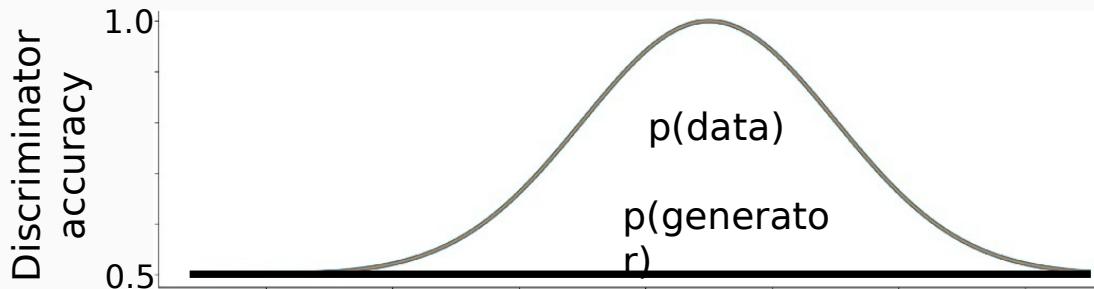
$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

vs

Discriminator

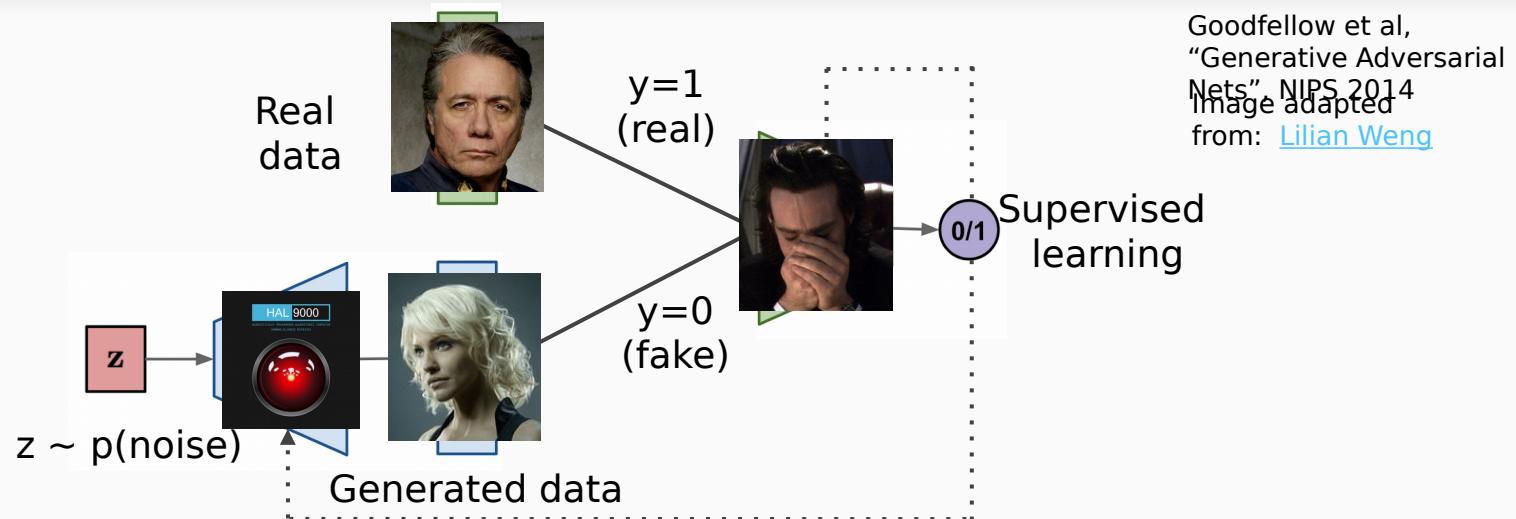
Tries to tell the difference between real examples (data) and fakes

$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$



Nash Equilibrium (Game theory)
 $p(\text{generator}) \approx p(\text{data})$

Implicit Density Estimation: Generative Adversarial Networks



Generator

Learns to transform noise into fakes in order to fool the Discriminator

$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

vs

Discriminator

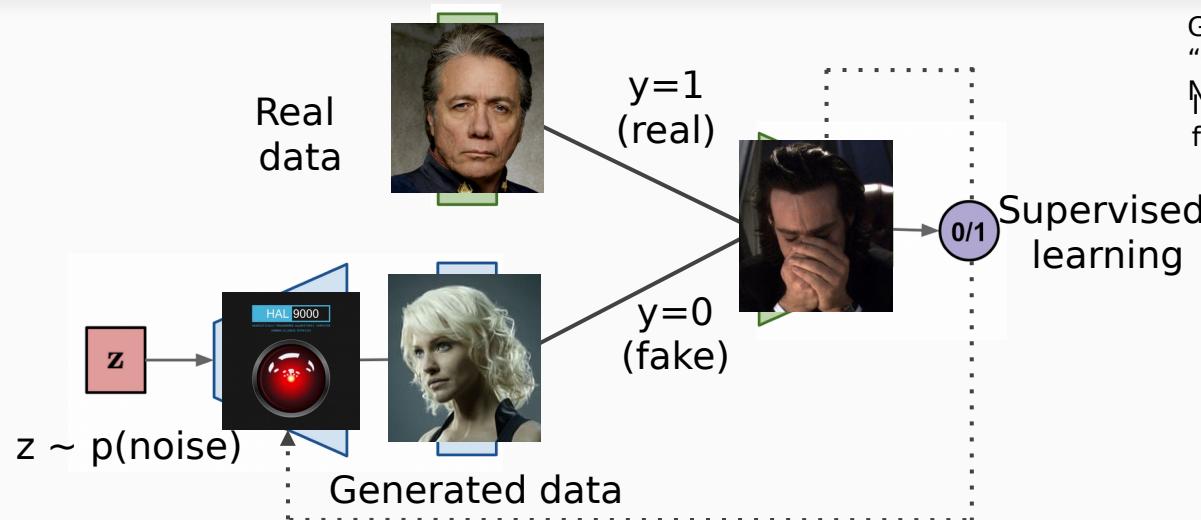
Tries to tell the difference between real examples (data) and fakes

$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

Rapid improvement in image quality over time



Implicit Density Estimation: Generative Adversarial Networks



Goodfellow et al,
"Generative Adversarial
Nets", NIPS 2014
Image adapted
from: [Lilian Weng](#)

Generator

Learns to transform noise into fakes in order to fool the Discriminator

$$\max_{\phi} \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(D_{\phi}(G_{\theta}(\mathbf{z})))]$$

vs

Discriminator

Tries to tell the difference between real examples (data) and fakes

$$\max_{\phi} \mathbb{E}_{\mathbf{x} \sim \mathbf{p}_{\text{data}}} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

Rapid improvement in image quality over time



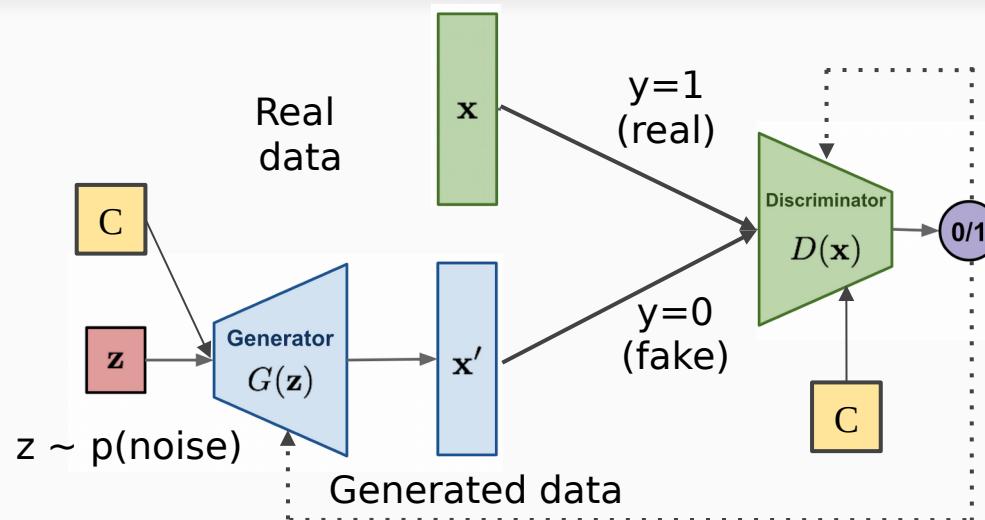
Many variants of GANs:

E.g. BiGAN, CycleGAN, BigGAN, InfoGAN, etc

See: [GAN Zoo](#) and

[Hong et al. ACM Computing Surveys, 2019](#)

Conditional GANs



E.g. text-image pairs
human)



Goodfellow et al,
“Generative Adversarial
Nets”, NIPS 2014
Image adapted
from: [Lilian Weng](#)

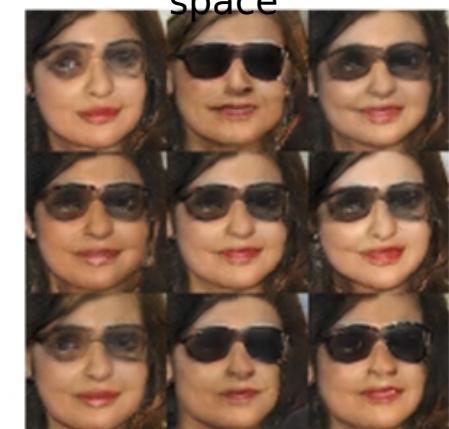
Supervised Generative Model

Latent space manipulations with GANs

Generated samples from latent vectors conditioned on text



Generated images from linearly manipulated latent space



Example from averaged latent vector



man
with glasses



man
without glasses



woman
without glasses



woman with glasses



Results of doing the same arithmetic in pixel space

Radford et al, ICLR 2016

Fun GAN examples

Image to image translation

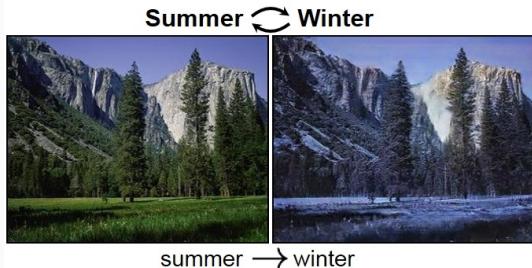


Pix2Pix using Conditional GANs (
[Isola et al. 2016](#))

Combine latent spaces
to create novel images



GANBreeder using BigGAN ([Brock et al. 2018](#))



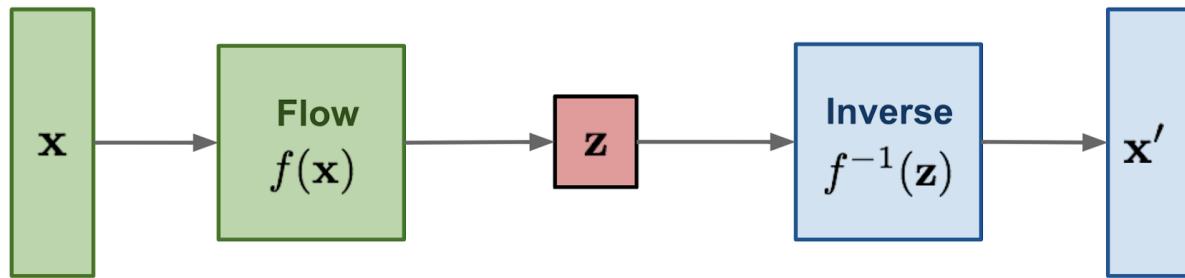
Style Transfer, Image to image
translation



Pix2Pix using
[CycleGAN](#)

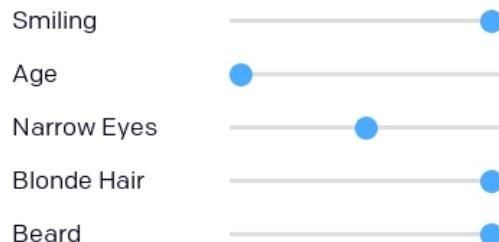
Flow-Based Models (Brief)

Flow-based generative models:
minimize the negative log-likelihood



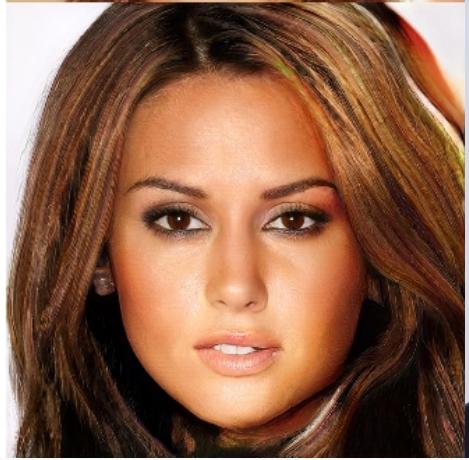
Explicitly learns probability density

Works in latent space



[Kingma, Dhariwal \(2018\) Glow: Generative Flow with Invertible 1x1 Convolutions](#)

Which reality TV star is fake real?



Which celebrity is fake real?

